

ESE345 Computer Architecture Logistics

Instructor: Ryan Thielke
E-mail: ryan.thielke@stonybrook.edu

Office Hours: M/W 4:00-5:45PM
Light Engineering Room 258A

Class: MW 6:05-7:25 PM
Earth & Space 131

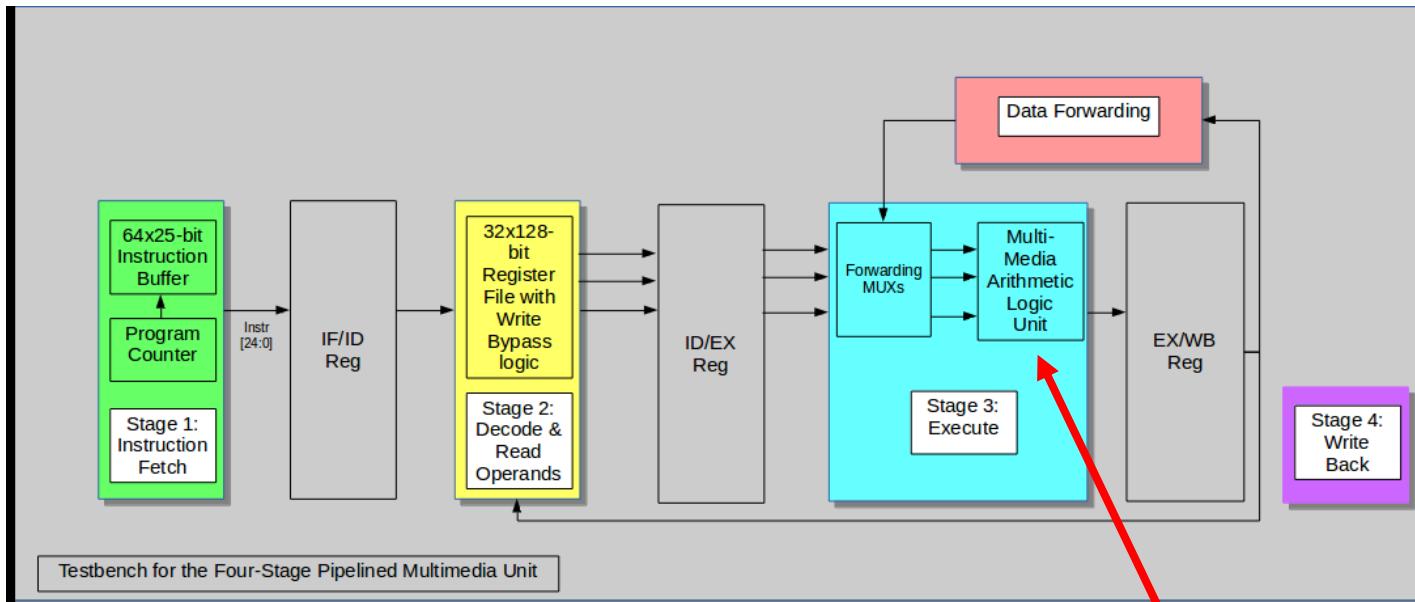
Teaching Assistants: Xiaokun Zhao Mondays 2-3pm (Project and Course Material)
Jingwei Li Tuesdays 2-3pm (HW and Course Material)

Text: David A. Patterson and John L. Hennessy “Computer Organization & Design The Hardware/Software Interface,” 5th ed., 2014 by Elsevier Inc. ISBN: 978-0-12-407726-3.

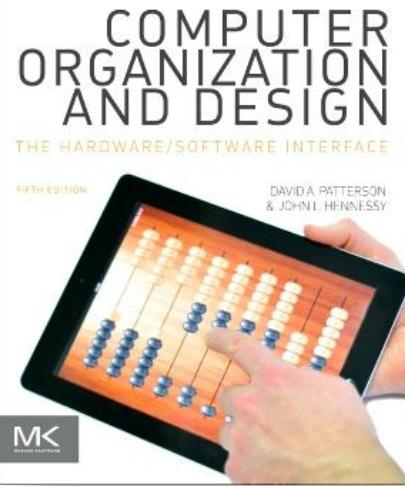
Course Grading: Six homeworks + 1 programming assignment : 14%
Exams (two midterms) : 64%
Project: 22%

Discord: <https://discord.gg/dS6qXUT9DU>

Project: VHDL Pipelined Multimedia Unit Model



- Read **Chapter 3.6** on subword parallelism to understand the original concept of multimedia processing
- Refresh your knowledge of VHDL/Verilog in the HDL design of digital circuits
- Develop and submit behavioral VHDL/Verilog code and its verification results for all multimedia ALU functions at the 3rd stage
- Develop the HDL model of the four-stage pipelined multimedia unit and its modules
- Verify individual modules of your design with their testbenches before instantiating them in higher order modules. Verify the final model with a testbench and show correct pipeline operation with and without forwarding



ESE 345 Computer Architecture

Introduction and Technology Trends

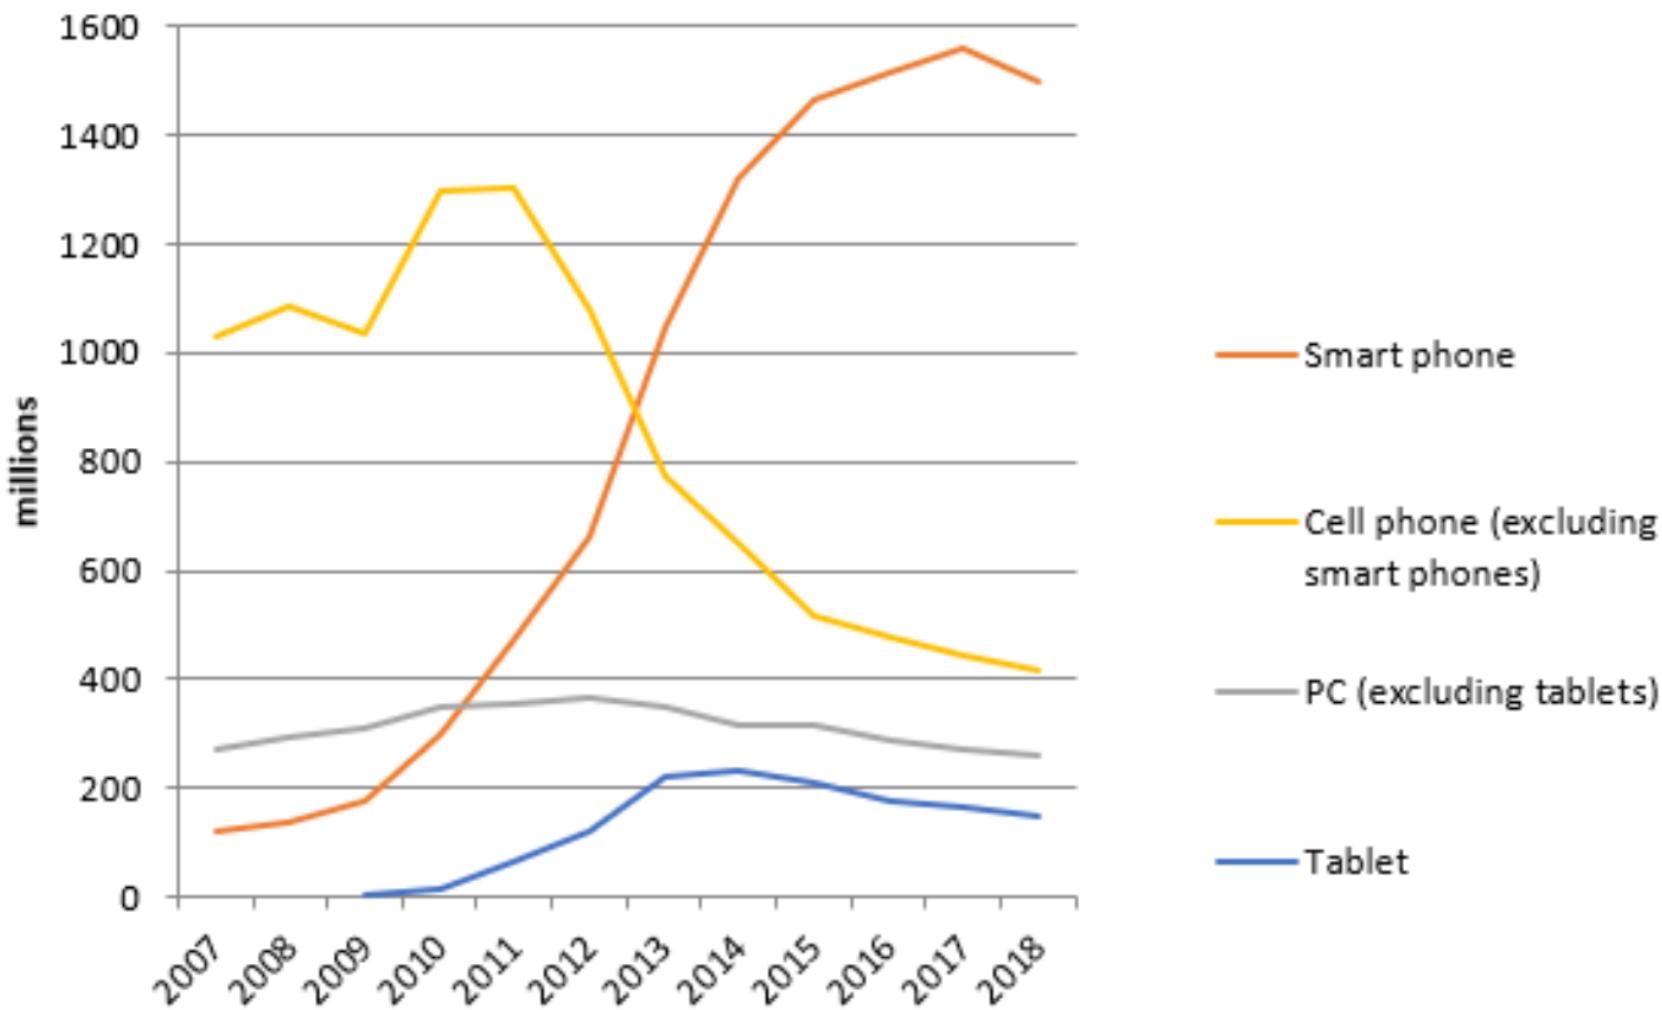
Classes of Computers

- Personal Computers
 - General purpose, variety of software
 - Subject to cost/performance tradeoff
- Server Computers
 - Network based
 - High capacity, performance, reliability
 - Range from small servers to building sized

Classes of Computers

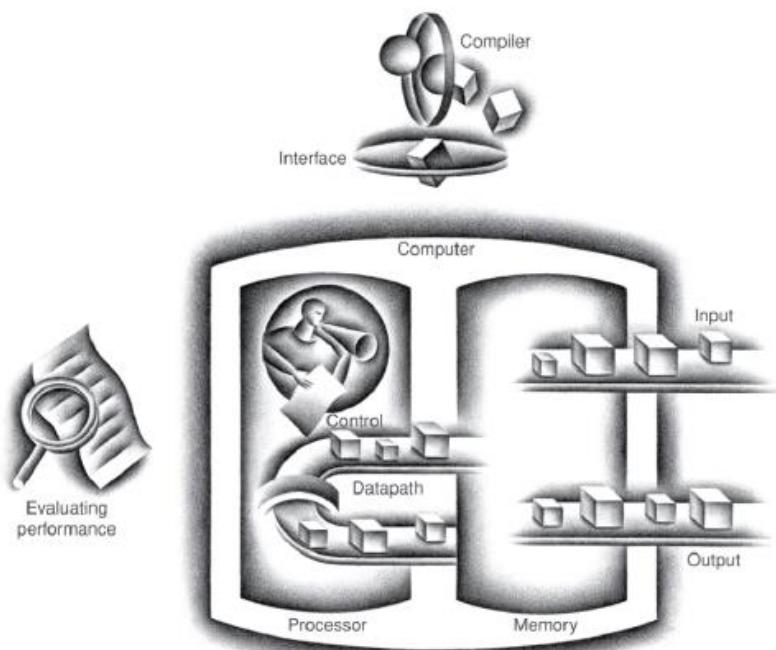
- Supercomputers
 - Type of server
 - High-end scientific and engineering calculations
 - Highest computing capability but represent a small fraction of the overall computer market
- Embedded Computers
 - Hidden as component of systems
 - Stringent power/performance/cost constraints

The PostPC Era



Components of a Computer

The BIG Picture



- Same components for all kinds of computers since 1946
 - Desktop, server, embedded
- Input/output includes
 - User-interface devices
 - Display, keyboard, mouse
 - Storage devices
 - Hard disk, CD/DVD, flash
 - Network adapters
 - For communicating with other computers

Inside the Processor Core (CPU)

- Datapath: performs operations on data
- Control: sequences datapath, memory, etc.
- Cache memory
 - Small fast SRAM memory for immediate access to data

What You Will Learn

- How programs are translated into the machine language
 - And how the hardware executes them
- The hardware/software interface
- What determines program performance
 - And how it can be improved
- How hardware designers improve performance
- What is parallel processing

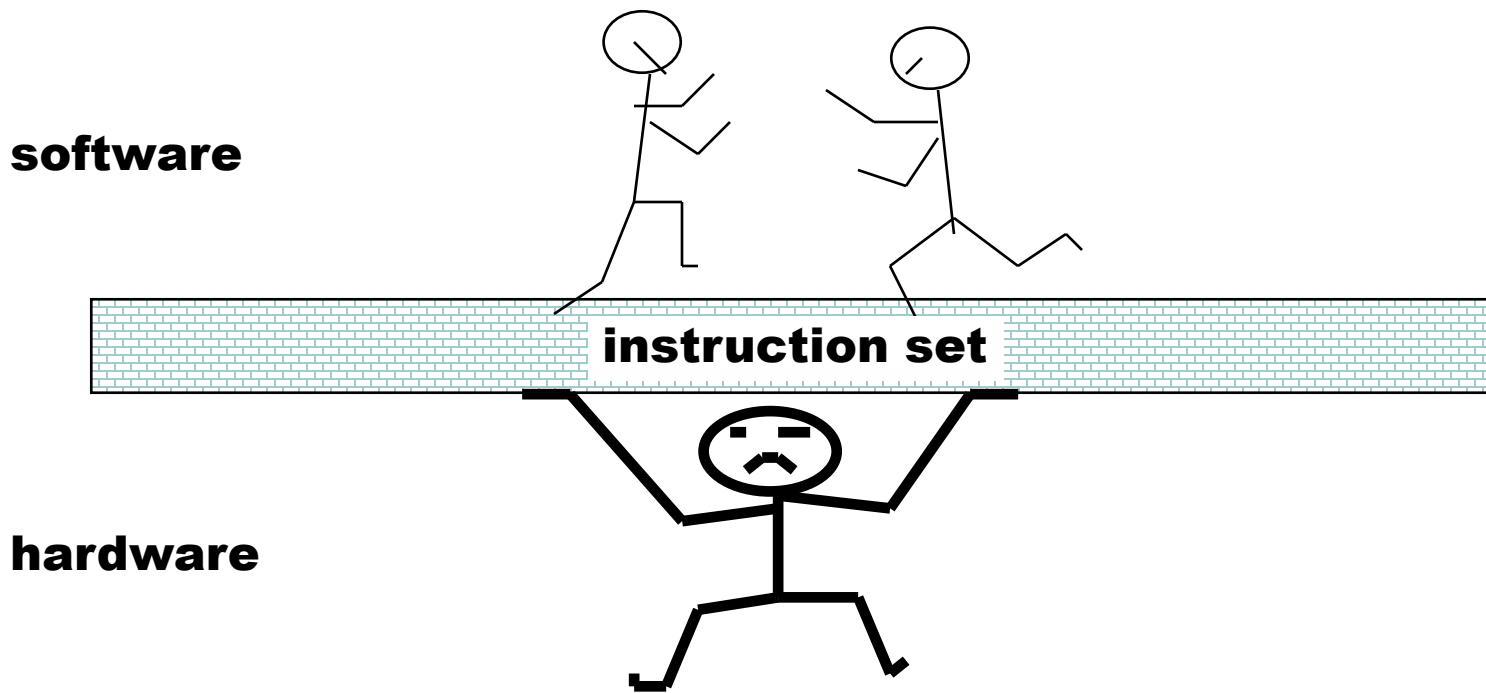
Understanding Performance

- Algorithm
 - Determines number of operations executed
- Programming language, compiler, architecture
 - Determine number of machine instructions executed per operation
- Processor and memory system
 - Determine how fast instructions are executed
- I/O system (including OS)
 - Determines how fast I/O operations are executed

What is Computer Architecture?

- Old definition of computer architecture
= instruction set architecture (ISA)
- Our view is computer architecture >> ISA
- Architect's job much more than instruction set design; technical hurdles today *more* challenging than those in instruction set design

Instruction Set Architecture: Critical Interface

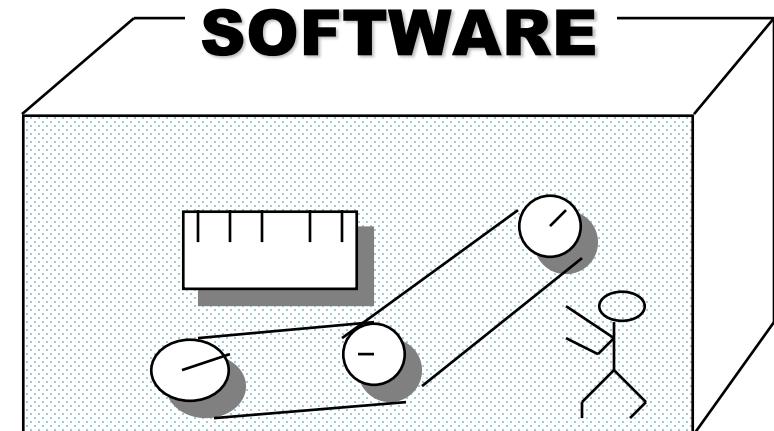


- Properties of a good abstraction
 - Lasts through many generations (portability)
 - Used in many different ways (generality)
 - Provides convenient functionality to higher levels
 - Permits an efficient implementation at lower levels

Instruction Set Architecture

“... the attributes of a [computing] system as seen by the programmer, *i.e.* the conceptual structure and functional behavior, as distinct from the organization of the data flows and controls the logic design, and the physical implementation.”
— Amdahl,
Blaauw, and Brooks, 1964

- **Organization of Programmable Storage**
- **Data Types & Data Structures:
Encodings & Representations**
- **Instruction Formats**
- **Instruction (or Operation Code) Set**
- **Modes of Addressing and Accessing Data Items and Instructions**
- **Exceptional Conditions**



Example: MIPS

r0	0
r1	
:	
:	
r31	
PC	
lo	
hi	

Programmable storage

$2^{32} \times$ bytes

31 x 32-bit GPRs (R0=0)

32 x 32-bit FP regs (paired DP)

HI, LO, PC

Data types ?

Format ?

**Addressing
Modes?**

Arithmetic logical

Add, AddU, Sub, SubU, And, Or, Xor, Nor, SLT, SLTU,
AddI, AddIU, SLTI, SLTIU, AndI, OrI, XorI, LUI
SLL, SRL, SRA, SLLV, SRLV, SRAV

Memory Access

LB, LBU, LH, LHU, LW, LWL, LWR
SB, SH, SW, SWL, SWR

Control

J, JAL, JR, JALR
BEq, BNE, BLEZ, BGTZ, BLTZ, BGEZ, BLTZAL, BGEZAL

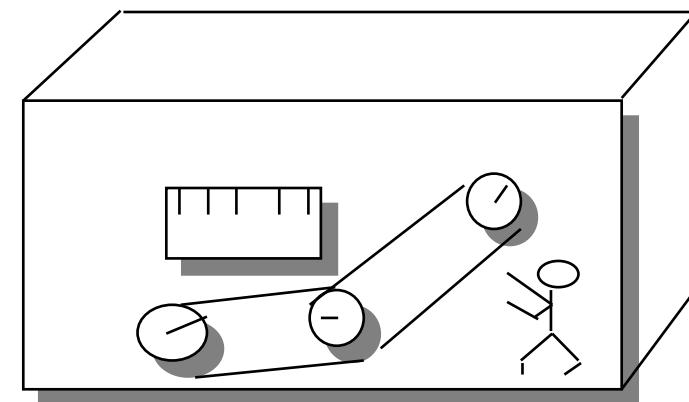
Organization (=Microarchitecture)

- Capabilities & Performance Characteristics of Principal Functional Units
 - (e.g., Registers, ALU, Shifters, Logic Units, ...)
- Ways in which these components are interconnected
- Information flows between components
- Logic and means by which such information flow is controlled.
- Choreography of FUs to realize the ISA
- Register Transfer Level (RTL) Description

Logic Designer's View

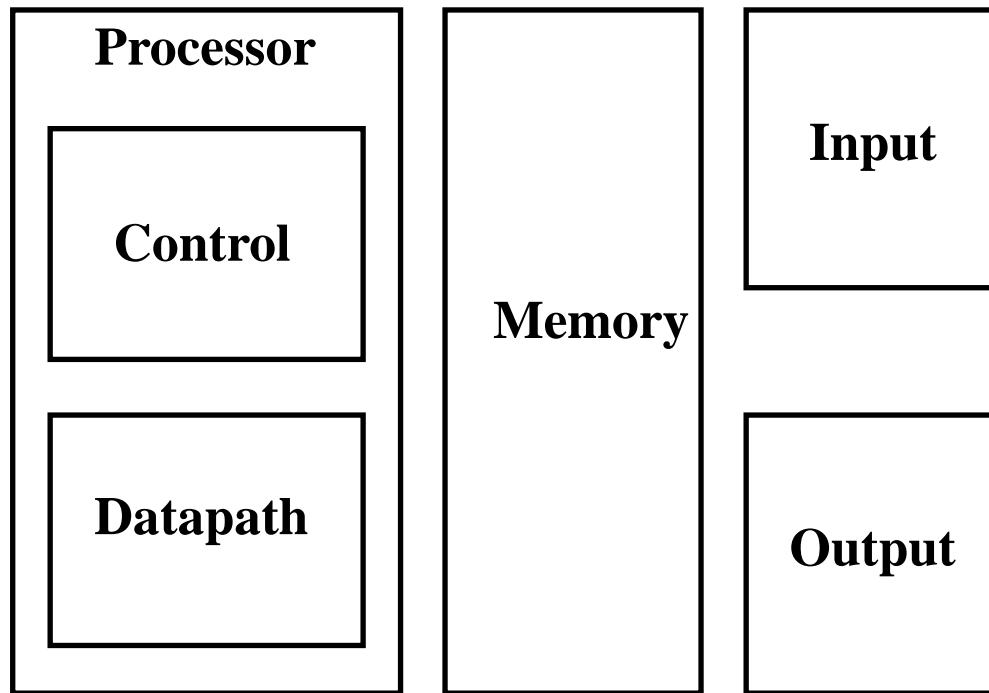
ISA Level

FUs & Interconnect

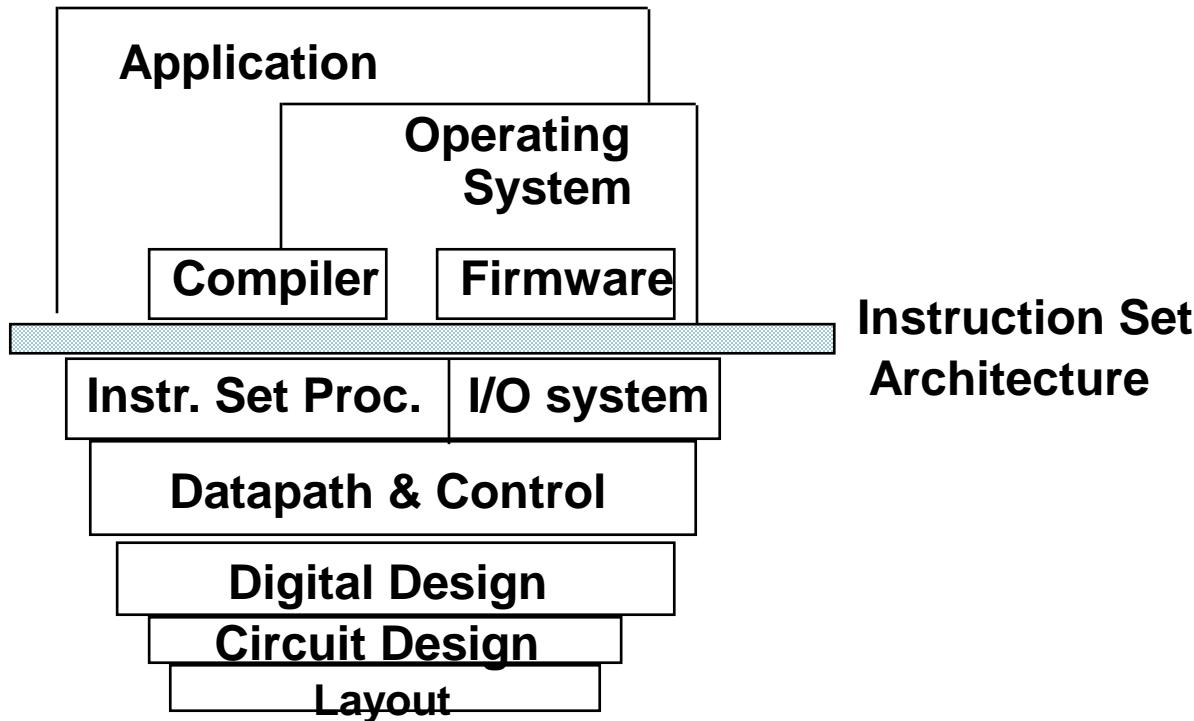


The Big Picture

- Since 1946 all computers have had 5 components



What is “Computer Architecture” Now?



- In its broad definition, computer architecture is the *design of the abstraction layers* to maximize performance within constraints (e.g., cost, power, and availability) that allow us to implement information processing applications efficiently using available manufacturing technologies.

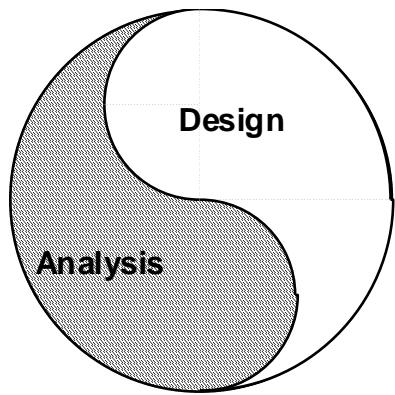
Great Ideas

- Use ***abstraction*** to simplify design
- Make the ***common case fast***
- Performance via ***parallelism***
- Performance via ***pipelining***
- Performance via ***prediction***
- ***Hierarchy*** of memories
- ***Dependability*** via redundancy
- ***Technology trends***

Computer Architecture is an Integrated Approach

- What really matters is the functioning of the complete system
 - hardware, runtime system, compiler, operating system, and application
 - In networking, this is called the “End to End argument”
- Computer architecture is not just about transistors, individual instructions, or particular implementations
 - E.g., Original RISC projects replaced complex instructions with a compiler + simple instructions

Computer Architecture is Design and Analysis

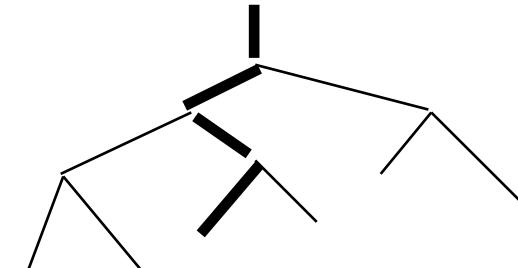


There is no single recipe for right architecture design

Architecture design is an iterative process:

- Searching the space of possible designs
- At all levels of computer systems

Creativity



Bad Ideas

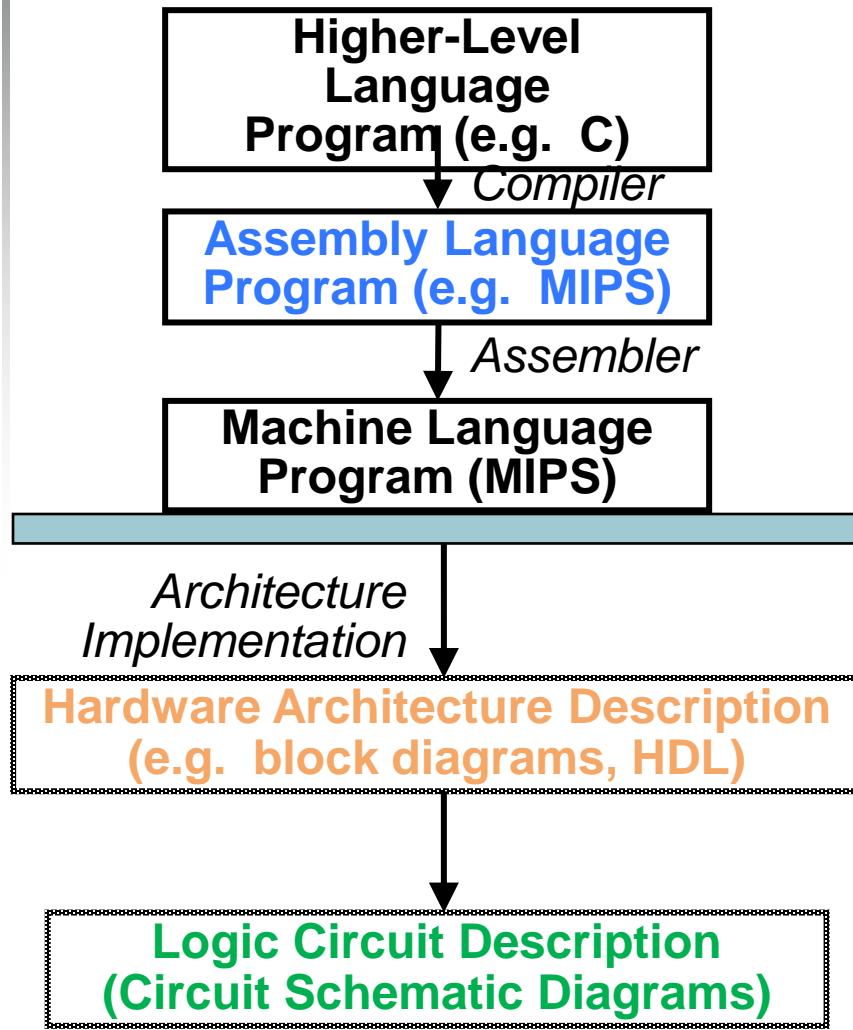
Mediocre Ideas

Good Ideas

Ideas in Computer Architecture

1. Levels of Representation
2. Technology Trends
3. Principle of Locality/Memory Hierarchy
4. Parallelism
5. Performance Measurement
6. Dependability via Redundancy

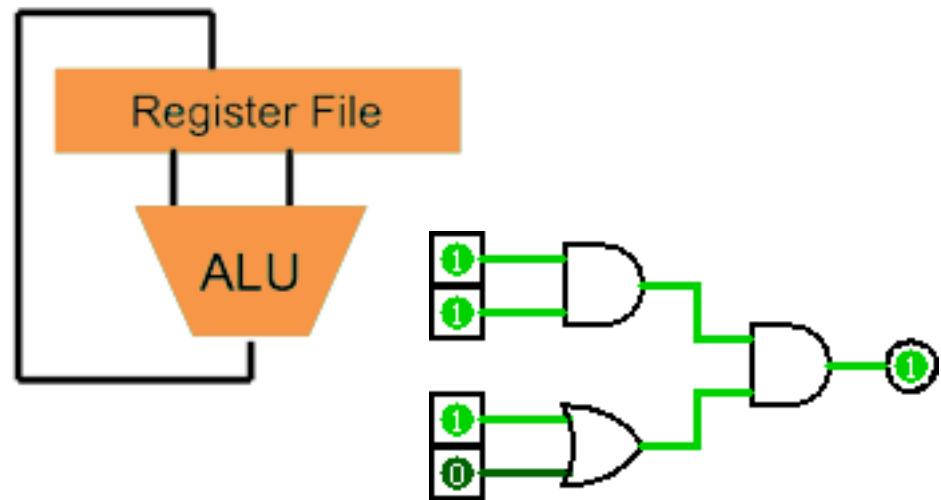
Idea #1: Levels of Representation



temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;

lw \$t0, 0(\$2)
lw \$t1, 4(\$2)
sw \$t1, 0(\$2)
sw \$t0, 4(\$2)

0000 1001 1100 0110 1010 1111 0101 1000
1010 1111 0101 1000 0000 1001 1100 0110
1100 0110 1010 1111 0101 1000 0000 1001
0101 1000 0000 1001 1100 0110 1010 1111



Levels of Program Code

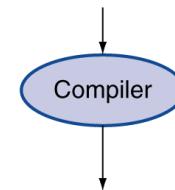
- High-level language
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
 - Assembly language
 - Textual representation of instructions
 - Hardware representation
 - Binary digits (bits)
 - Encoded instructions and data

High-level
language
program
(in C)

```

swap(int v[], int k)
{int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}

```

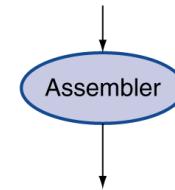


Assembly language program (for MIPS)

```

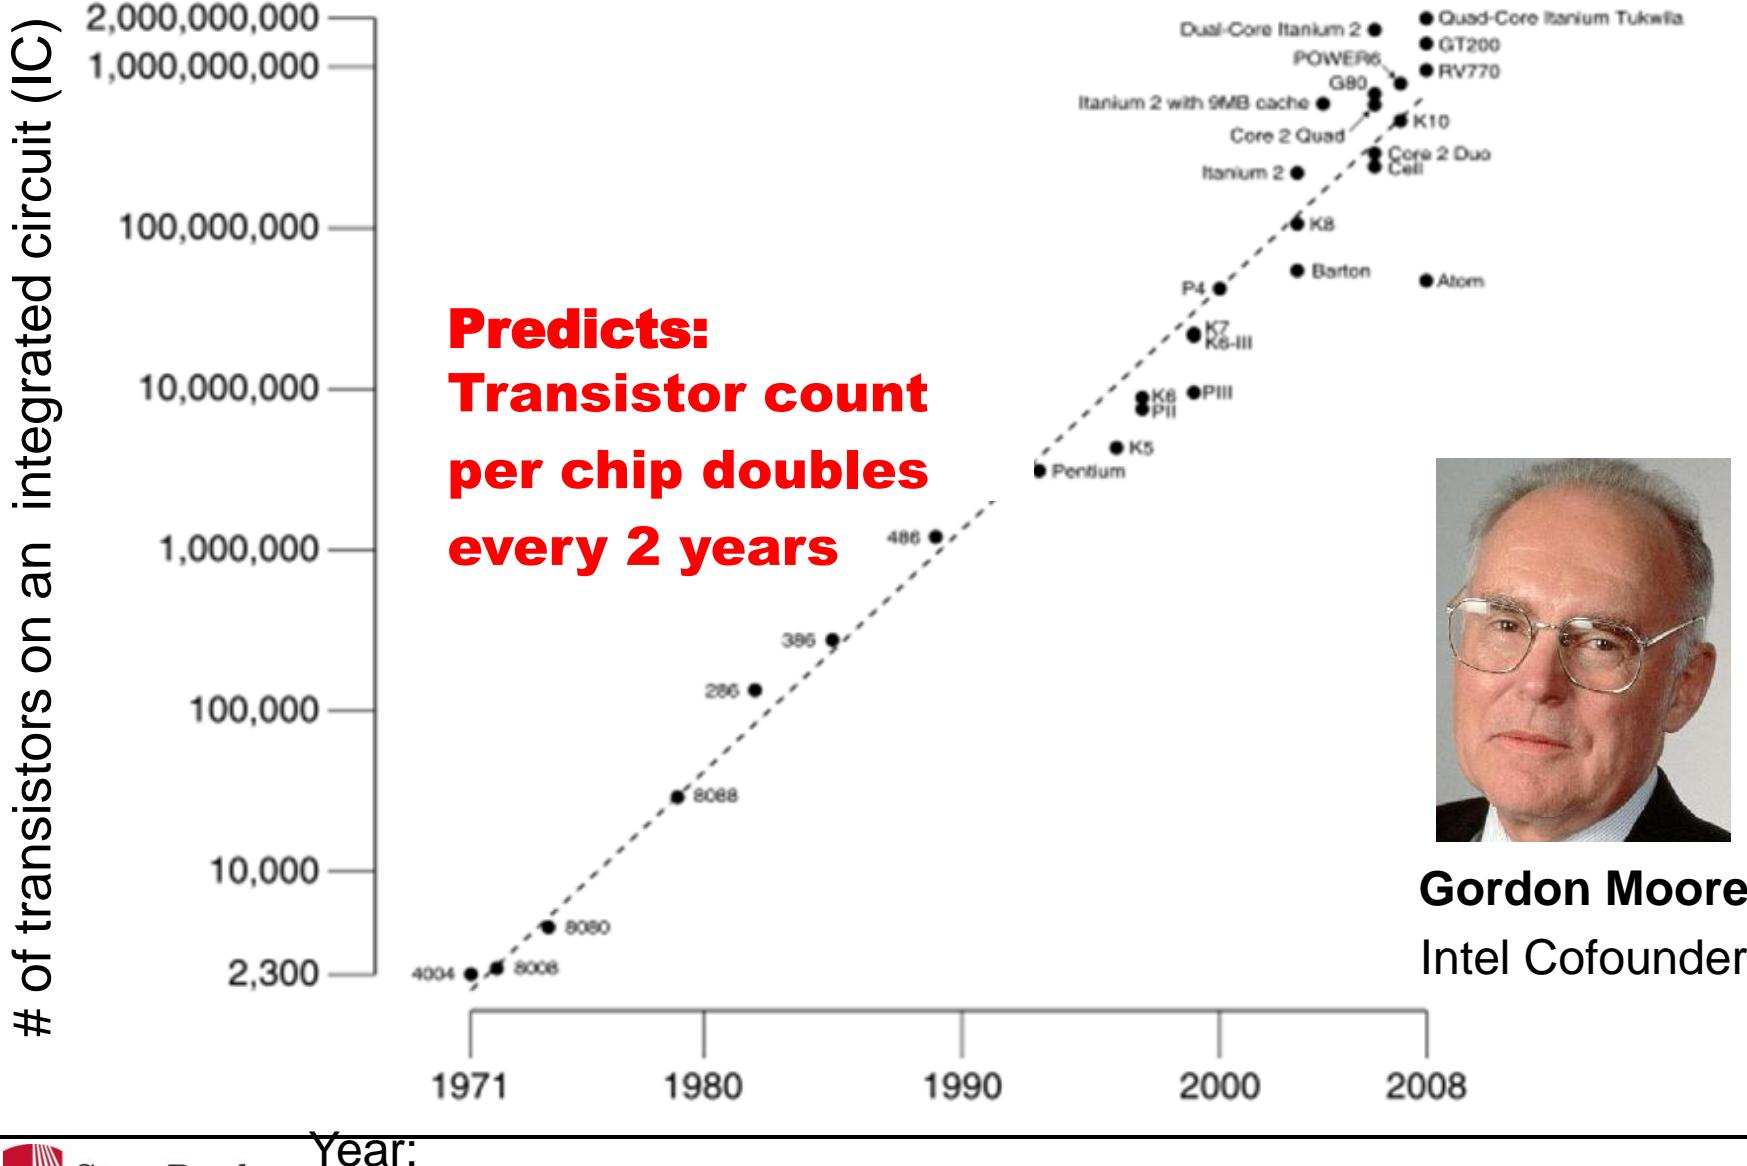
swap:
    muli $2, $5,4
    add  $2, $4,$2
    lw   $15, 0($2)
    lw   $16, 4($2)
    sw   $16, 0($2)
    sw   $15, 4($2)
    jr   $31

```



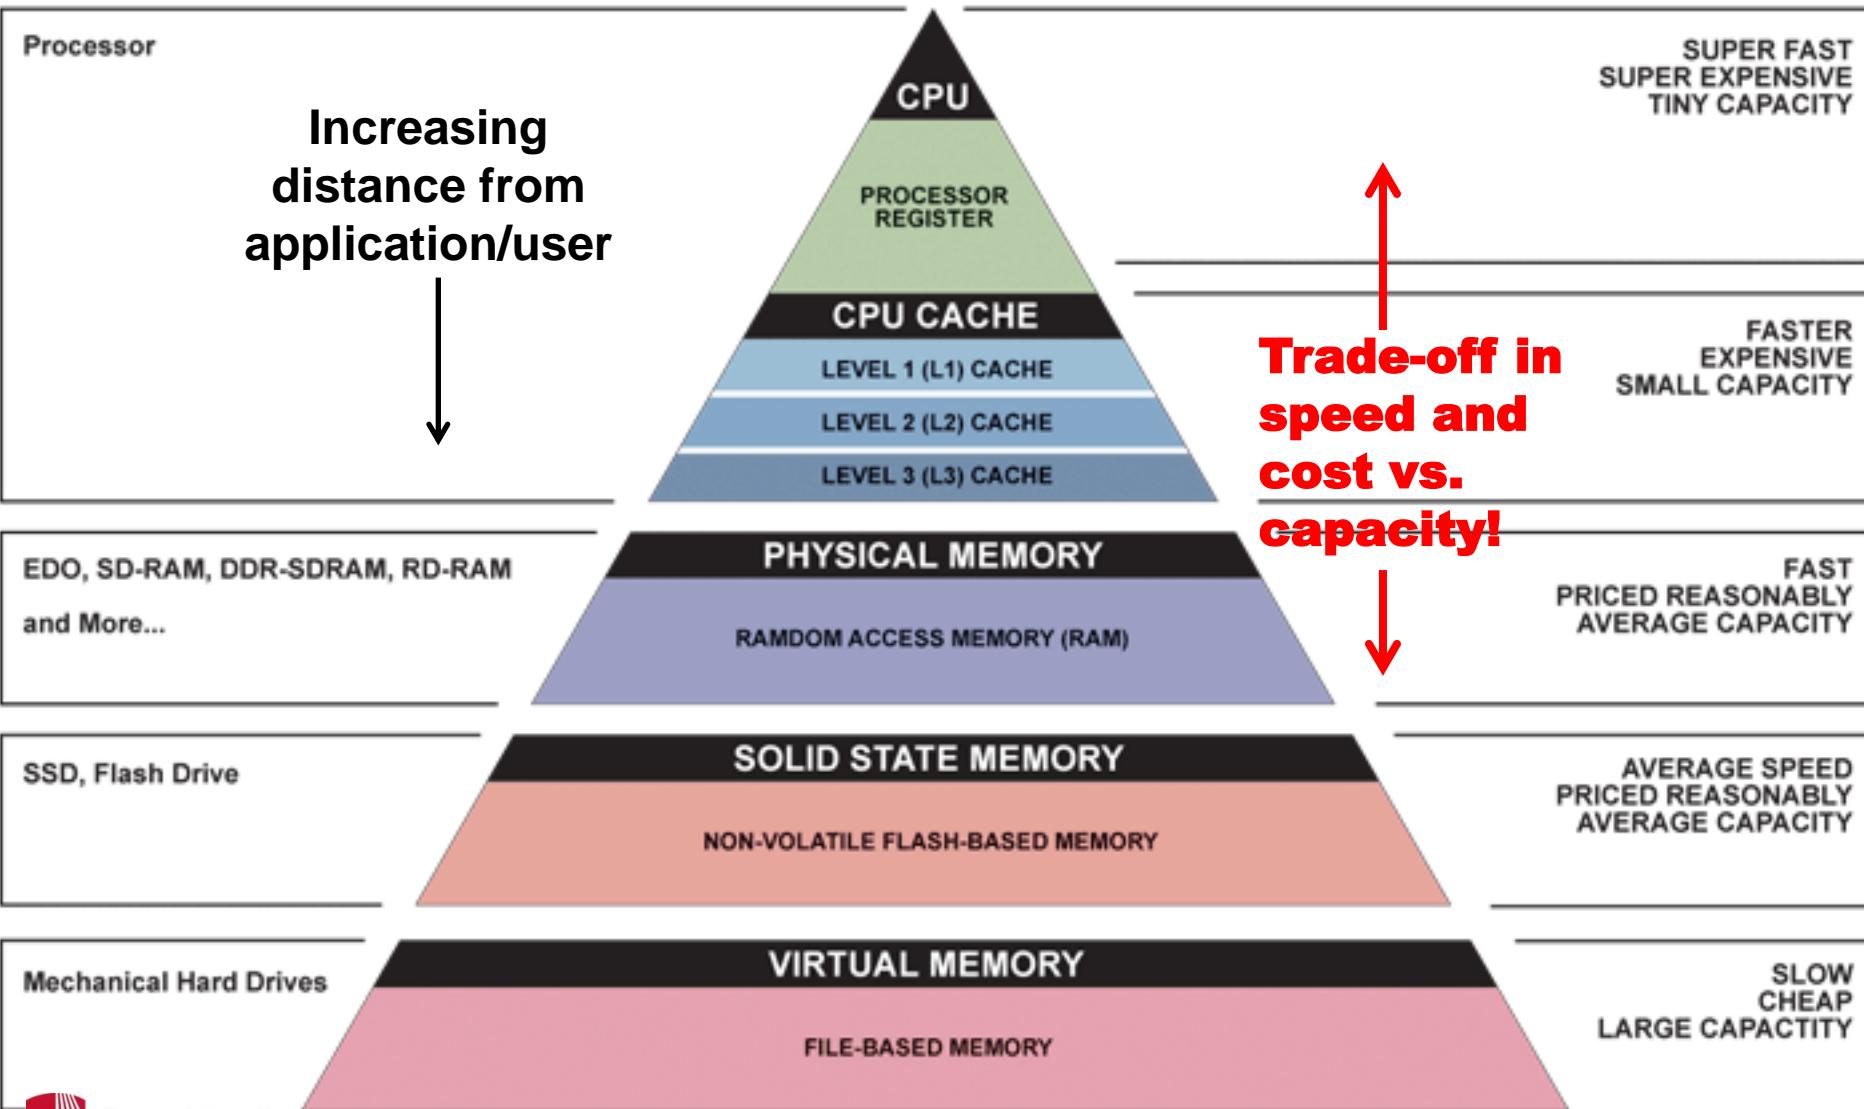
Binary machine language program (for MIPS)

Idea #2: Technology Trends

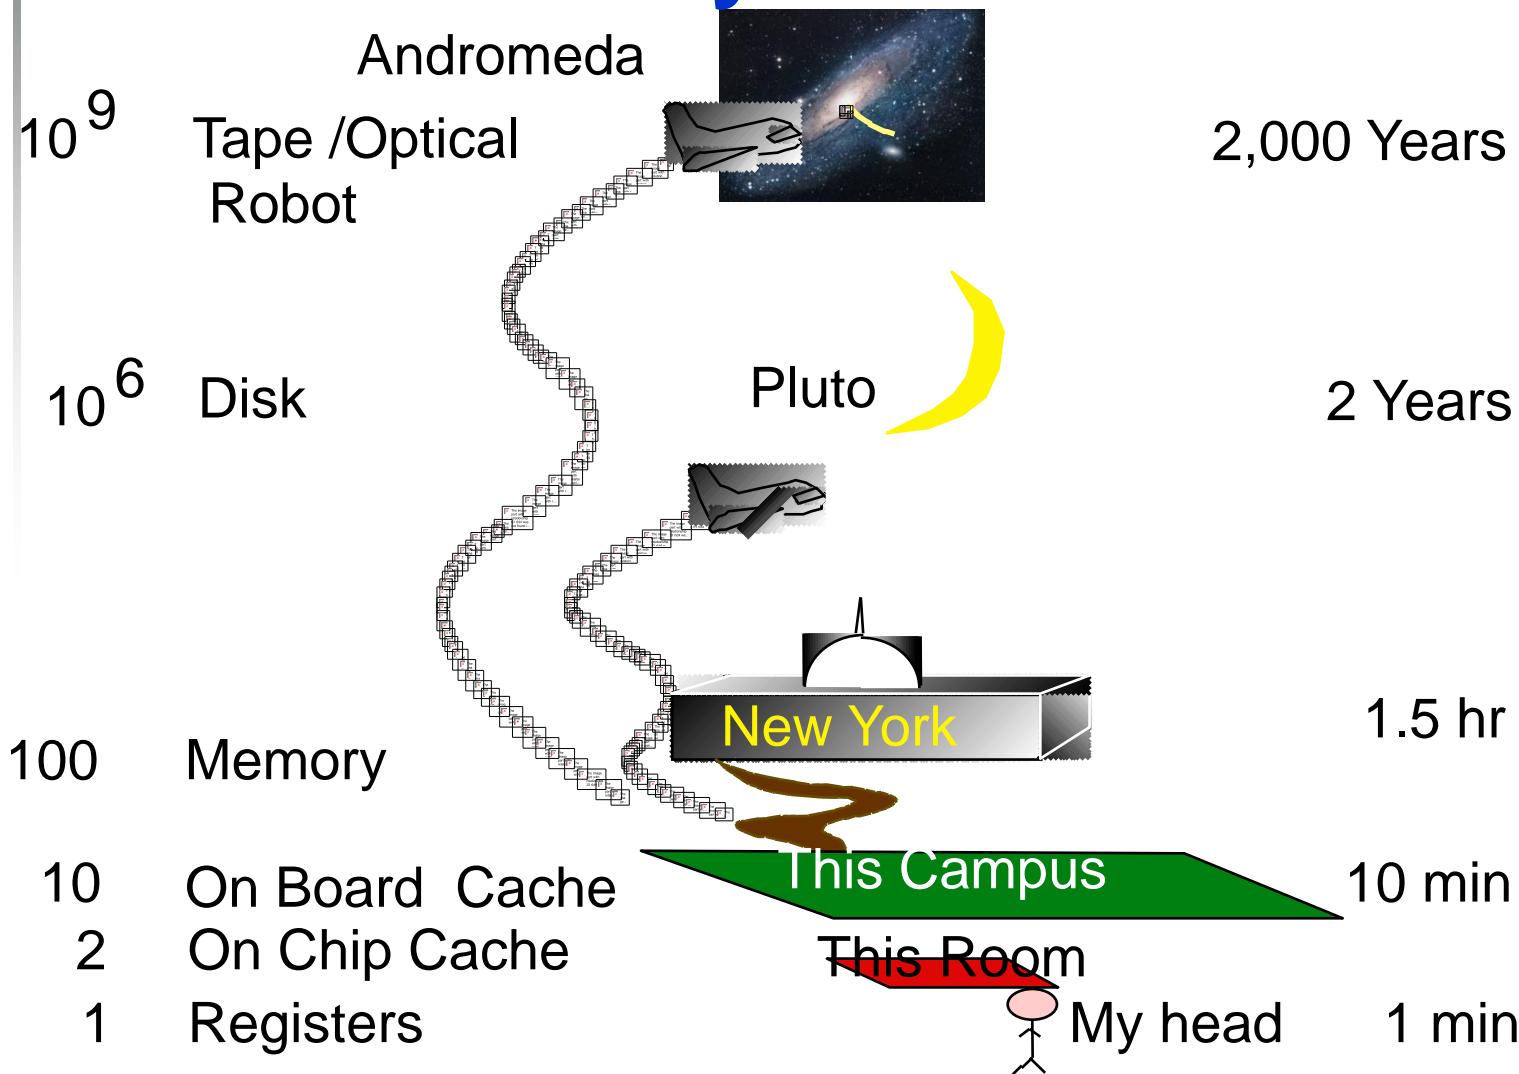


Gordon Moore
Intel Cofounder

Idea #3: Principle of Locality/ Memory Hierarchy



Storage Latency Analogy: How Far Away is the Data?



Memory Technology Costs

- Static RAM (SRAM)
 - 0.5ns – 2.5ns, \$500 – \$1000 per GB
- Dynamic RAM (DRAM)
 - 50ns – 70ns, \$3 – \$6 per GB
- Magnetic disk
 - 5ms – 20ms, \$0.01 – \$0.02 per GB
- Ideal memory
 - Access time of SRAM
 - Capacity and cost/GB of disk

Idea #4: Parallelism

- **Parallel Requests**
Assigned to computer
e.g. search “Garcia”
- **Parallel Threads**
Assigned to core
e.g. lookup, ads
- **Parallel Instructions**
> 1 instruction @ one time
e.g. 5 pipelined instructions
- **Parallel Data**
> 1 data item @ one time
e.g. add of 4 pairs of words
- **Hardware descriptions**
All gates functioning in parallel at same time

Hardware

Warehouse Scale Computer



Smart Phone



Leverage Parallelism & Achieve High Performance



Computer

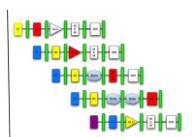
Core ... Core

Memory

Input/Output

Core

Instruction Unit(s)

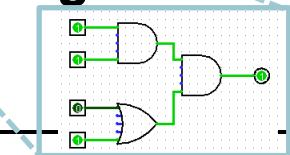


Functional Unit(s)

$A_0+B_0 A_1+B_1 A_2+B_2 A_3+B_3$

Cache Memory

Logic Gates

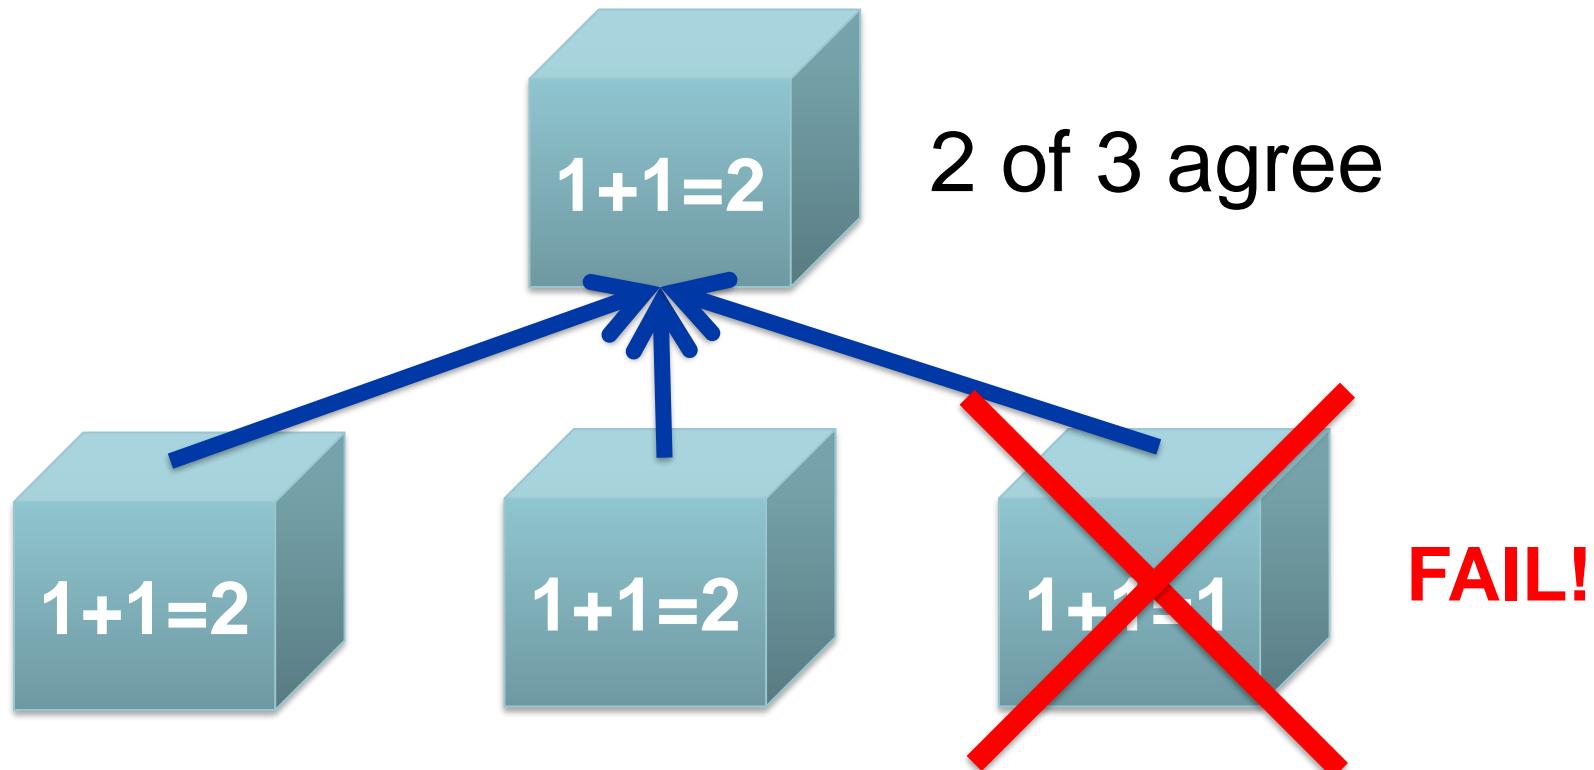


Idea #5: Performance Measurement and Improvement

- Allows direct comparisons of architectures and quantification of improvements
- Most common measures are *time to finish* (latency) and *rate of execution* (throughput)
- Match application and hardware to exploit:
 - Locality, parallelism, special hardware features

Idea #6: Dependability via Redundancy

- Redundancy so that a failing piece doesn't make the whole system fail



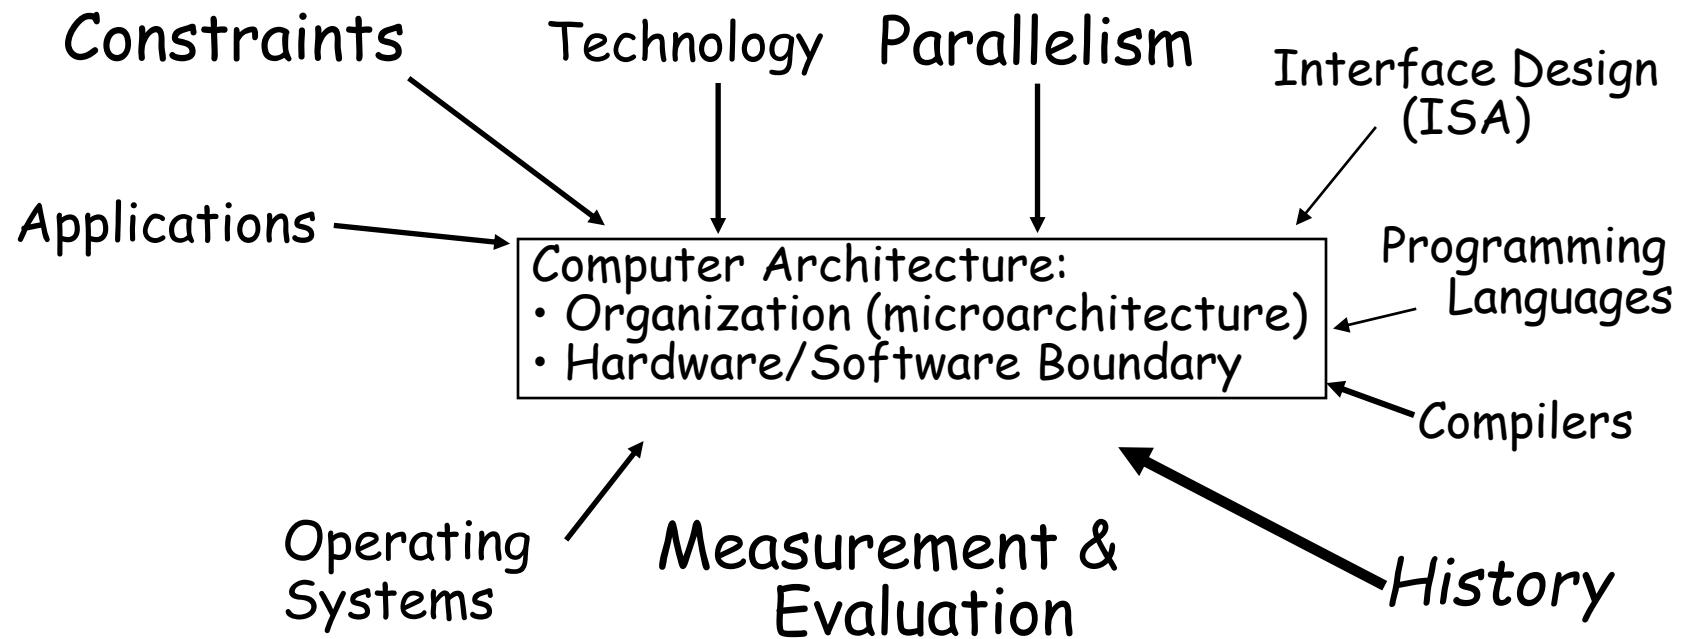
Idea #6: Dependability via Redundancy

- Applies to everything from datacenters to storage to memory
 - Redundant datacenters so that can lose 1 datacenter but Internet service stays online
 - Redundant disks so that can lose 1 disk but not lose data (Redundant Arrays of Independent Disks/RAID)
 - Redundant memory bits so that can lose 1 bit but no data (Error Correcting Code/ECC Memory)
- Increasing transistor density reduces the cost of redundancy



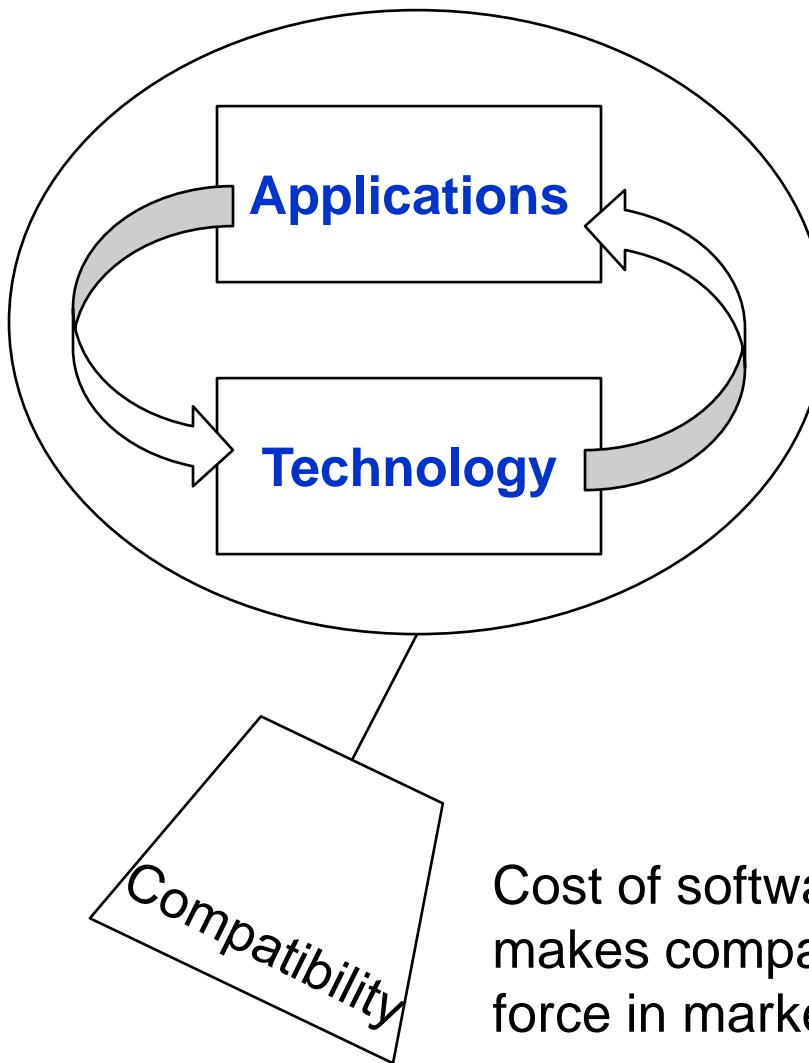
Computer Architecture Course Focus

Understanding the design techniques, machine structures, technology factors, evaluation methods that will determine the form of computers in 21st Century



Architecture Continually Changing

Applications suggest how to improve technology, provide revenue to fund development



Improved technologies make new applications possible

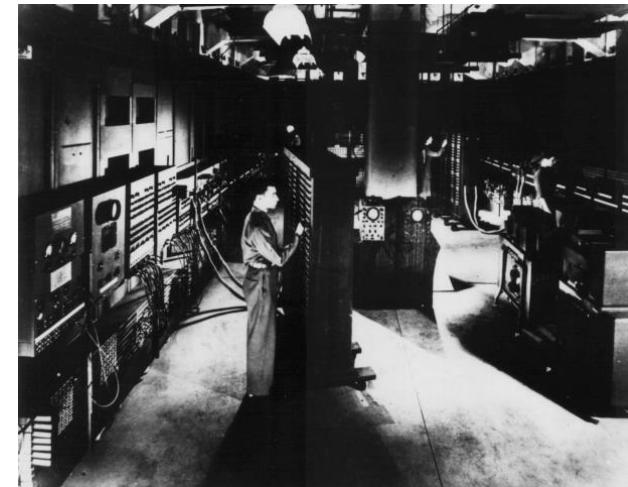
Cost of software development makes compatibility a major force in market

Electronic Numerical Integrator and Computer (ENIAC)

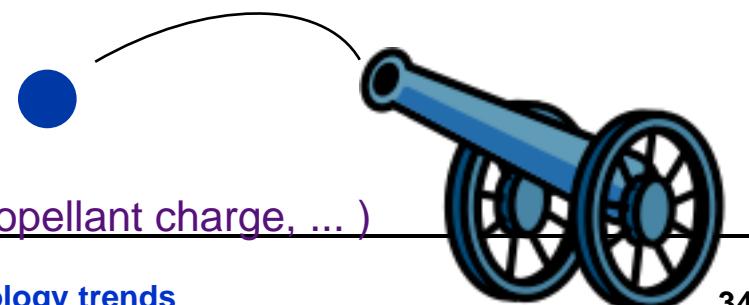
- Inspired by Atanasoff and Berry, Eckert and Mauchly designed and built ENIAC (1943-45) at the University of Pennsylvania
- The first, completely electronic, operational, general-purpose analytical calculator!
 - 30 tons, 72 square meters, 200KW
 - Each of the twenty 10 digit registers was 2 feet long
 - Used 18,000 vacuum tubes
- Performance
 - Read in 120 cards per minute
 - Addition took 200 ms, Division 6 ms
 - 1000 times faster than Mark I
- Not very reliable!

Application: Ballistic calculations

angle = f (location, tail wind, cross wind,
air density, temperature, weight of shell, propellant charge, ...)



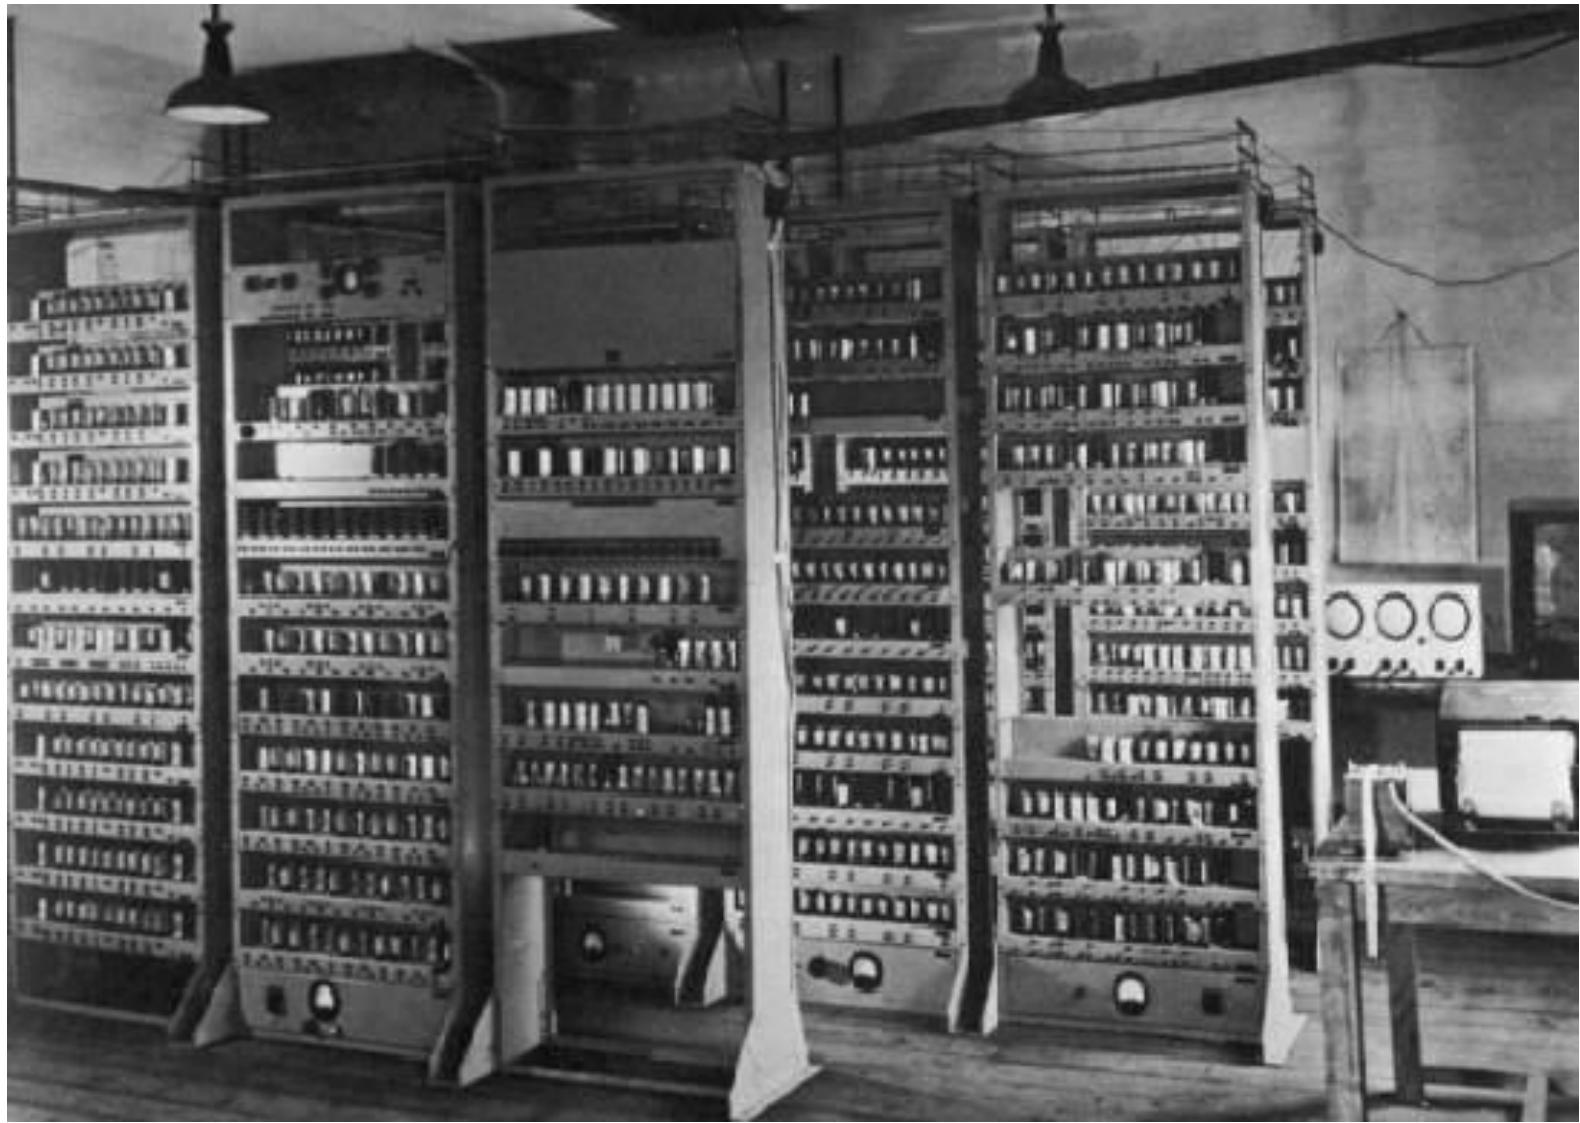
The slogan was: The ENIAC is able to calculate the trajectory of a large-caliber naval shell in less time than the shell takes to reach its target!



Electronic Discrete Variable Automatic Computer (EDVAC)

- ENIAC's programming system was external
 - Sequences of instructions were executed independently of the results of the calculation
 - Human intervention required to take instructions "out of order"
- Eckert, Mauchly, John von Neumann and others designed EDVAC (1944) to solve this problem
 - Solution was the *stored program computer*
⇒ "*program can be manipulated as data*"
- *First Draft of a report on EDVAC* was published in 1945, but just had von Neumann's signature!
 - In 1973 the court of Minneapolis attributed the honor of *inventing the computer* to John Atanasoff

EDSAC, University of Cambridge, UK, 1949



Dominant Problem: *Reliability*

Mean time between failures (MTBF)

MIT's Whirlwind with an MTBF of 20 min. was perhaps the most reliable machine !

Reasons for unreliability:

1. Vacuum Tubes
2. Storage medium
 - acoustic delay lines
 - mercury delay lines
 - Williams tubes
 - Selections

Reliability solved by invention of Core memory by J. Forrester 1954 at MIT for Whirlwind project

Technology Issues

ENIAC

18,000 tubes

20 10-digit numbers

⇒

EDVAC

4,000 tubes

2000 word storage

mercury delay lines

*ENIAC had many asynchronous parallel units
but only one was active at a time*

BINAC : Two processors that checked each other
for reliability.

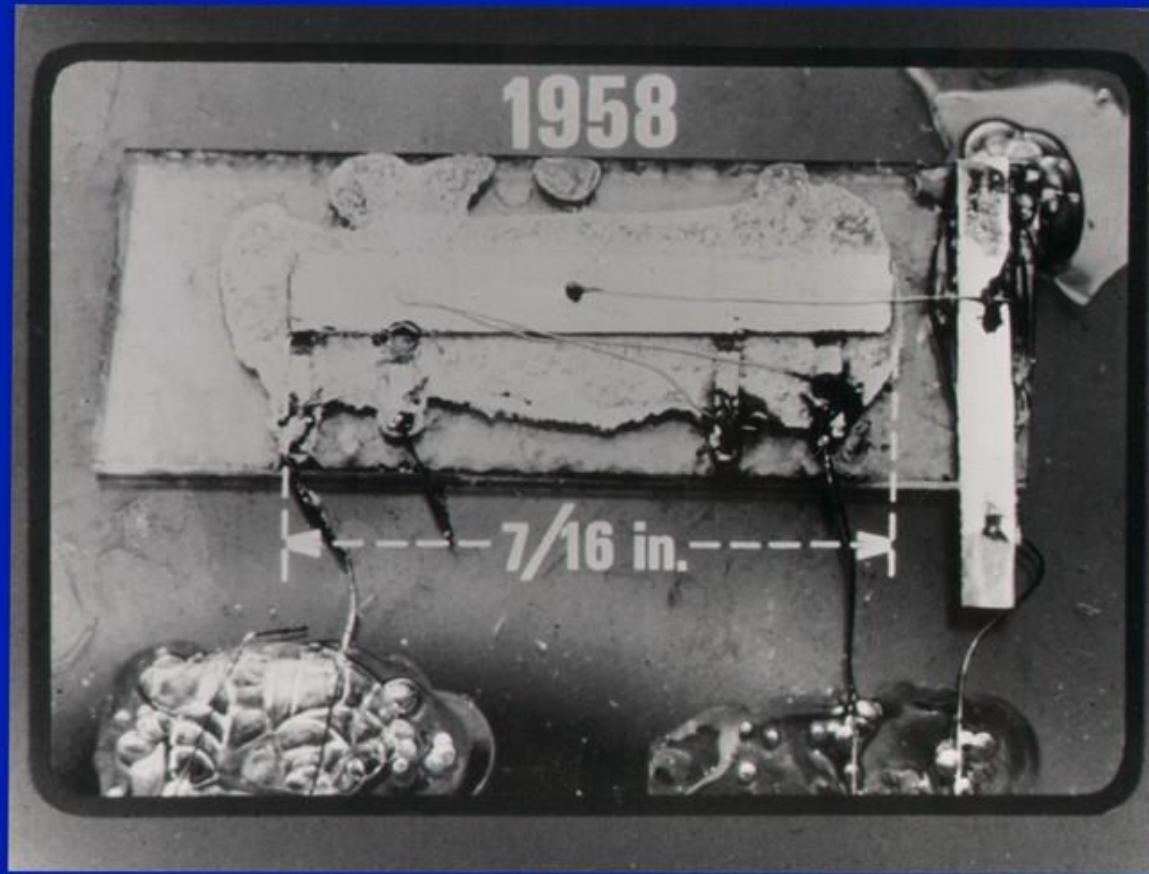
*Didn't work well because processors never
agreed*

Computers in mid 50's

- Hardware was expensive
- Stores were small (1000 words)
 - ⇒ No resident system software!
- Memory access time was 10 to 50 times slower than the processor cycle
 - ⇒ Instruction execution time was totally dominated by the *memory reference time*.
- The *ability to design complex control circuits* to execute an instruction was the central design concern as opposed to *the speed* of decoding or an ALU operation
- Programmer's view of the machine was inseparable from the actual hardware implementation
 - No software abstraction layer

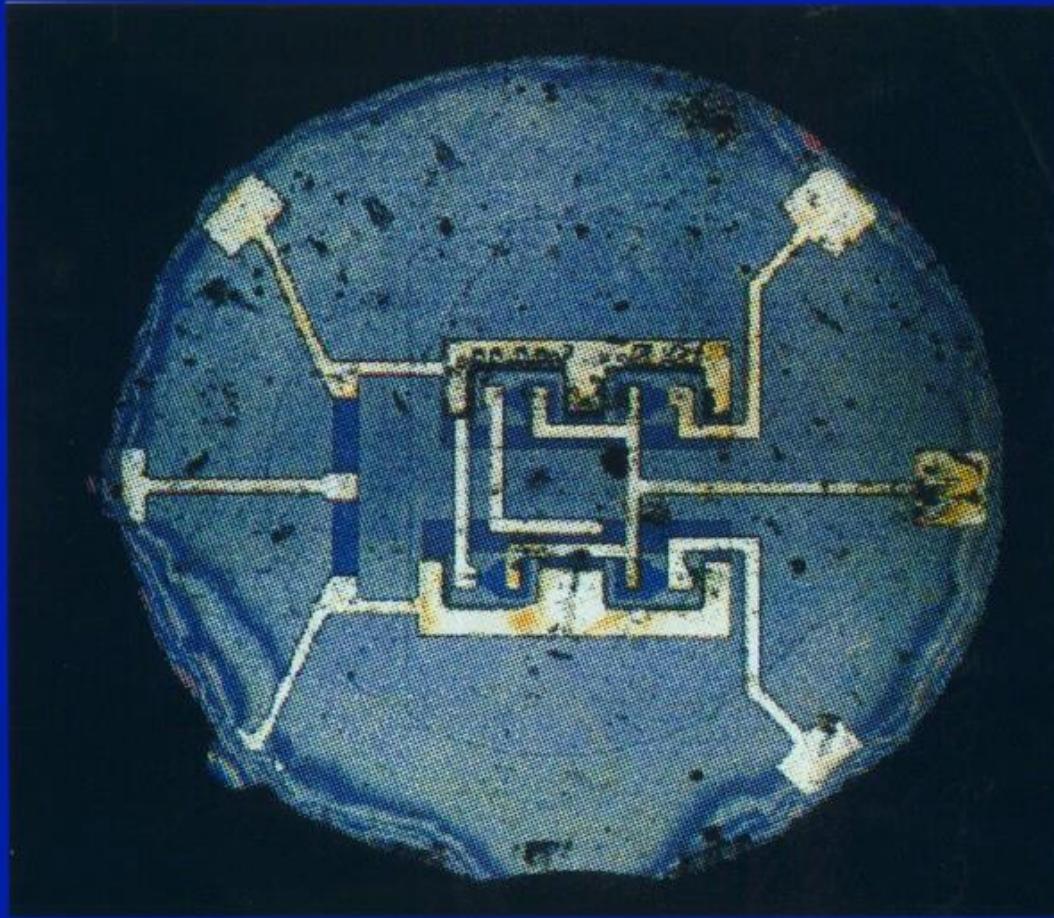
FIRST INTEGRATED CIRCUIT BY J. S. KILBY

(US Patent 3,138,763 filed Feb. 1959, granted 1964)

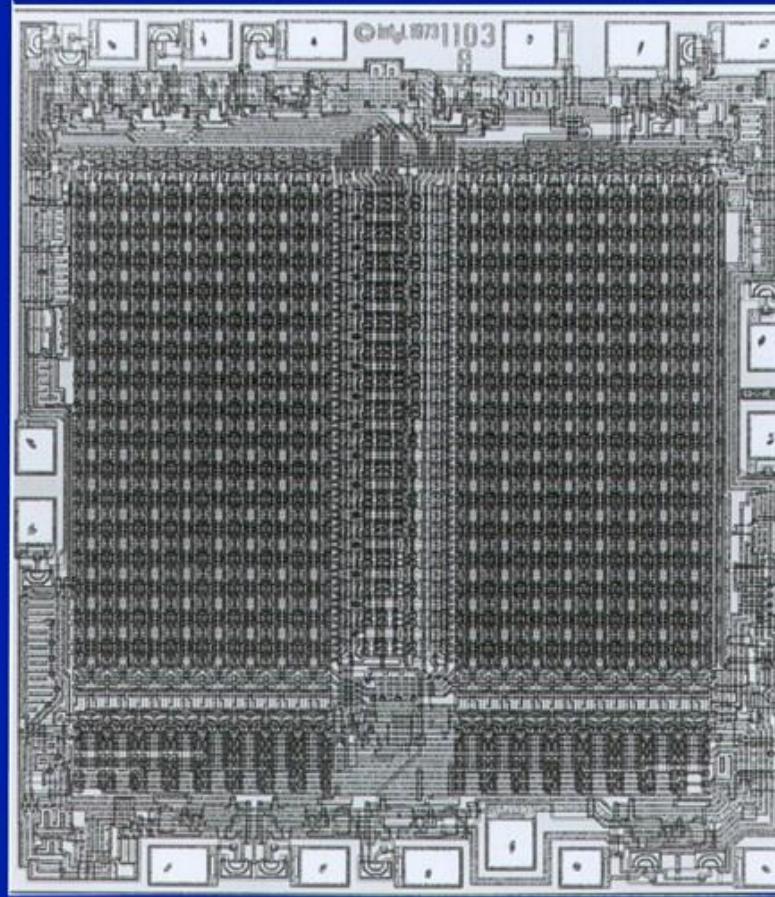


FIRST MONOLITHIC IC BY R. N. NOYCE

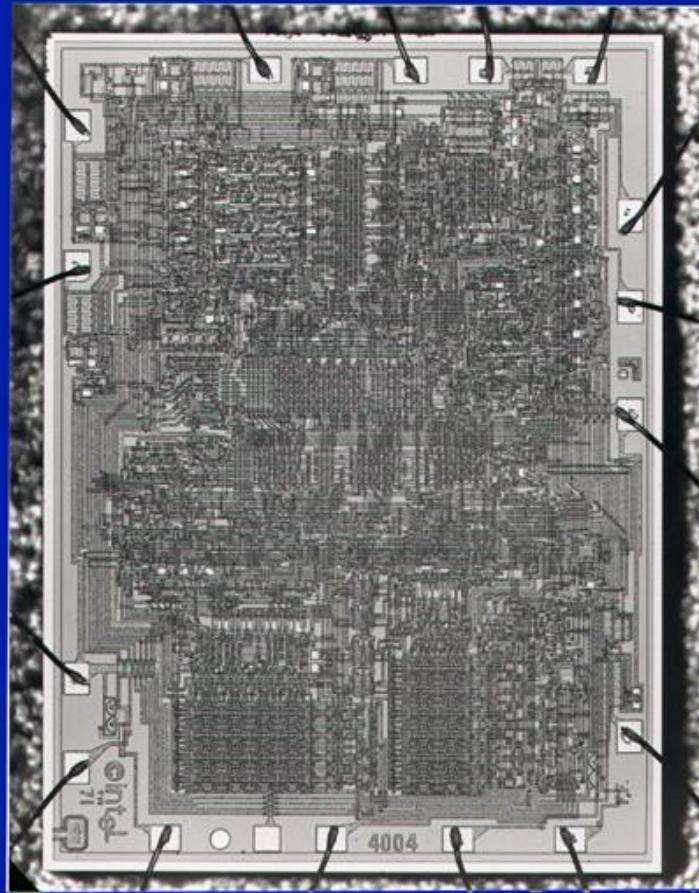
(US Patent 2,981,877 filed July 1959, granted 1961)



FIRST DRAM (1103 by Intel 1970)

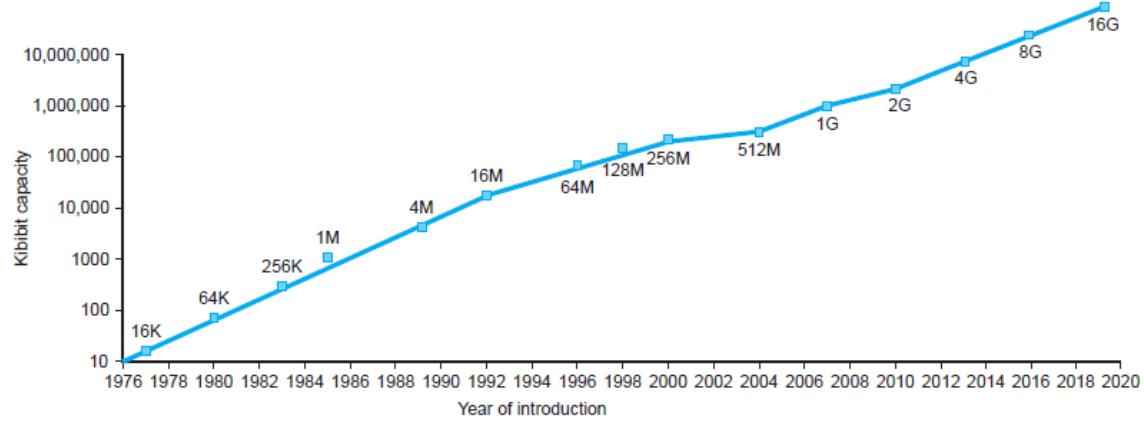


FIRST MICROPROCESSOR (4004 by Intel 1971)



Technology Trends

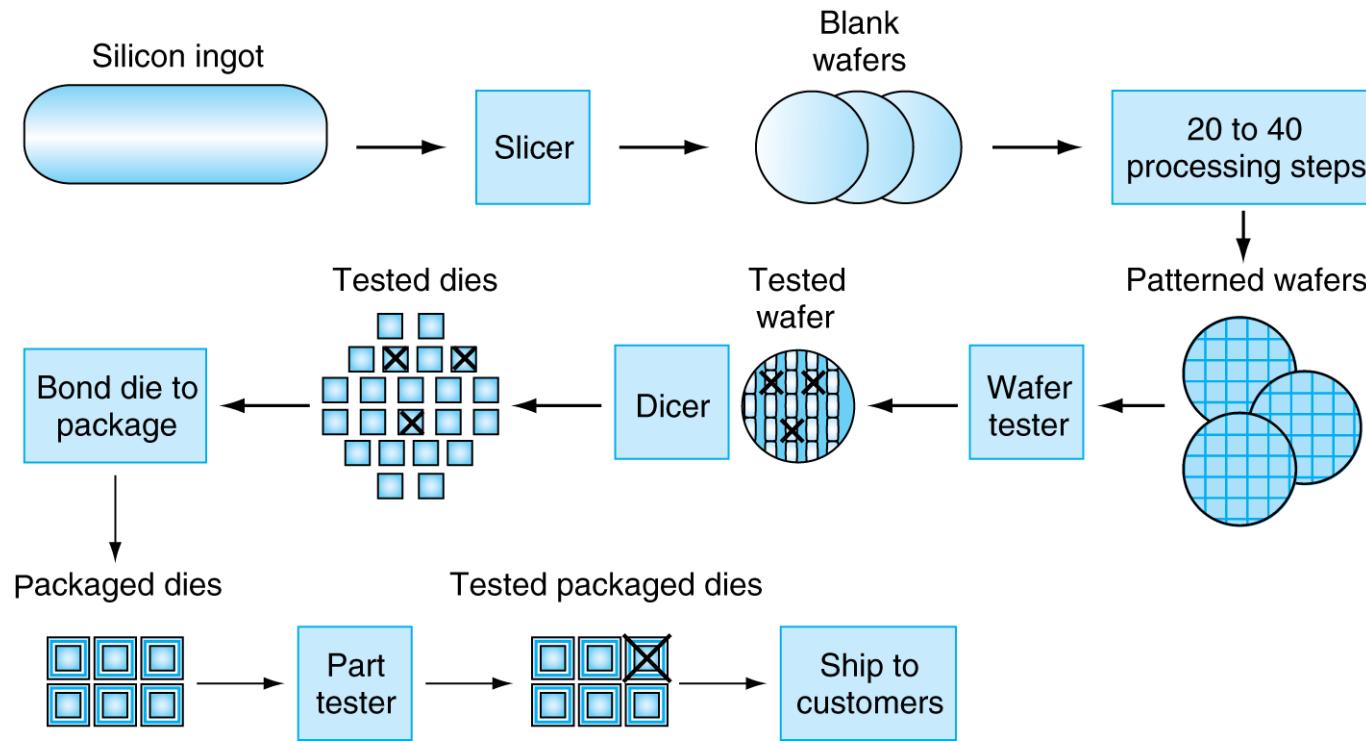
- Electronics technology continues to evolve
 - Increased capacity and performance
 - Reduced cost



DRAM capacity

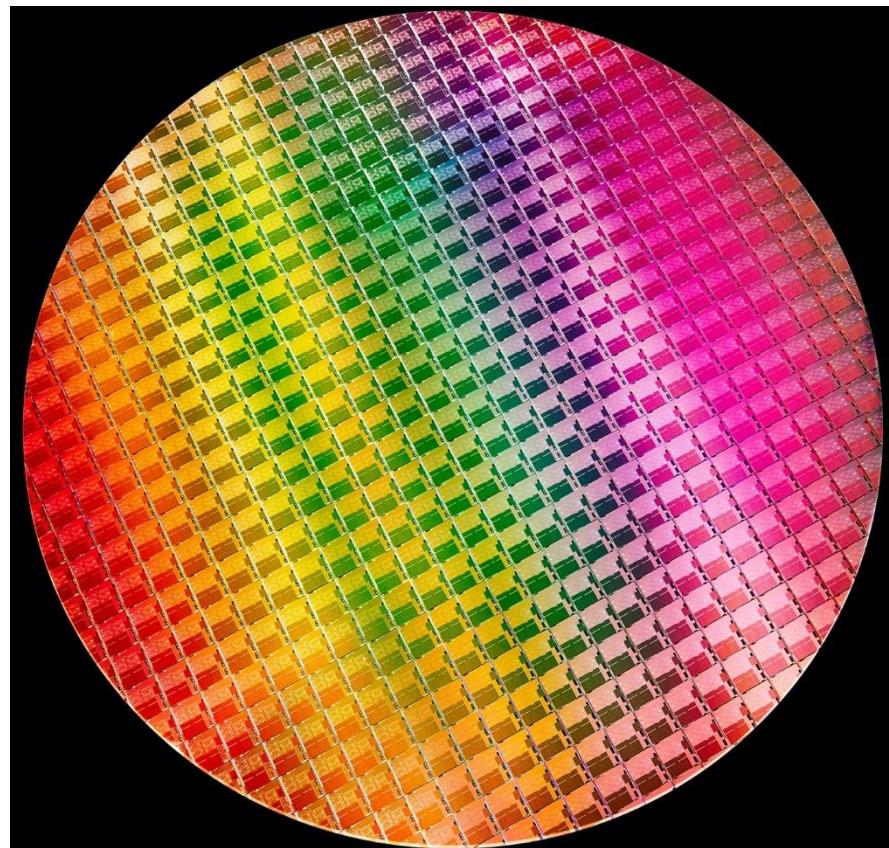
Year	Technology	Relative performance/cost
1951	Vacuum tube	1
1965	Transistor	35
1975	Integrated circuit (IC)	900
1995	Very large scale IC (VLSI)	2,400,000
2020	Ultra large scale IC	500,000,000,000

Manufacturing ICs



- Yield: proportion of working dies per wafer

Intel® Core 10th Generation



- 300mm wafer, 506 chips, 10nm technology
- Each chip is 11.4 x 10.7 mm

Integrated Circuit Cost

$$\text{Cost per die} = \frac{\text{Cost per wafer}}{\text{Dies per wafer} \times \text{Yield}}$$

$$\text{Dies per wafer} \approx \text{Wafer area}/\text{Die area}$$

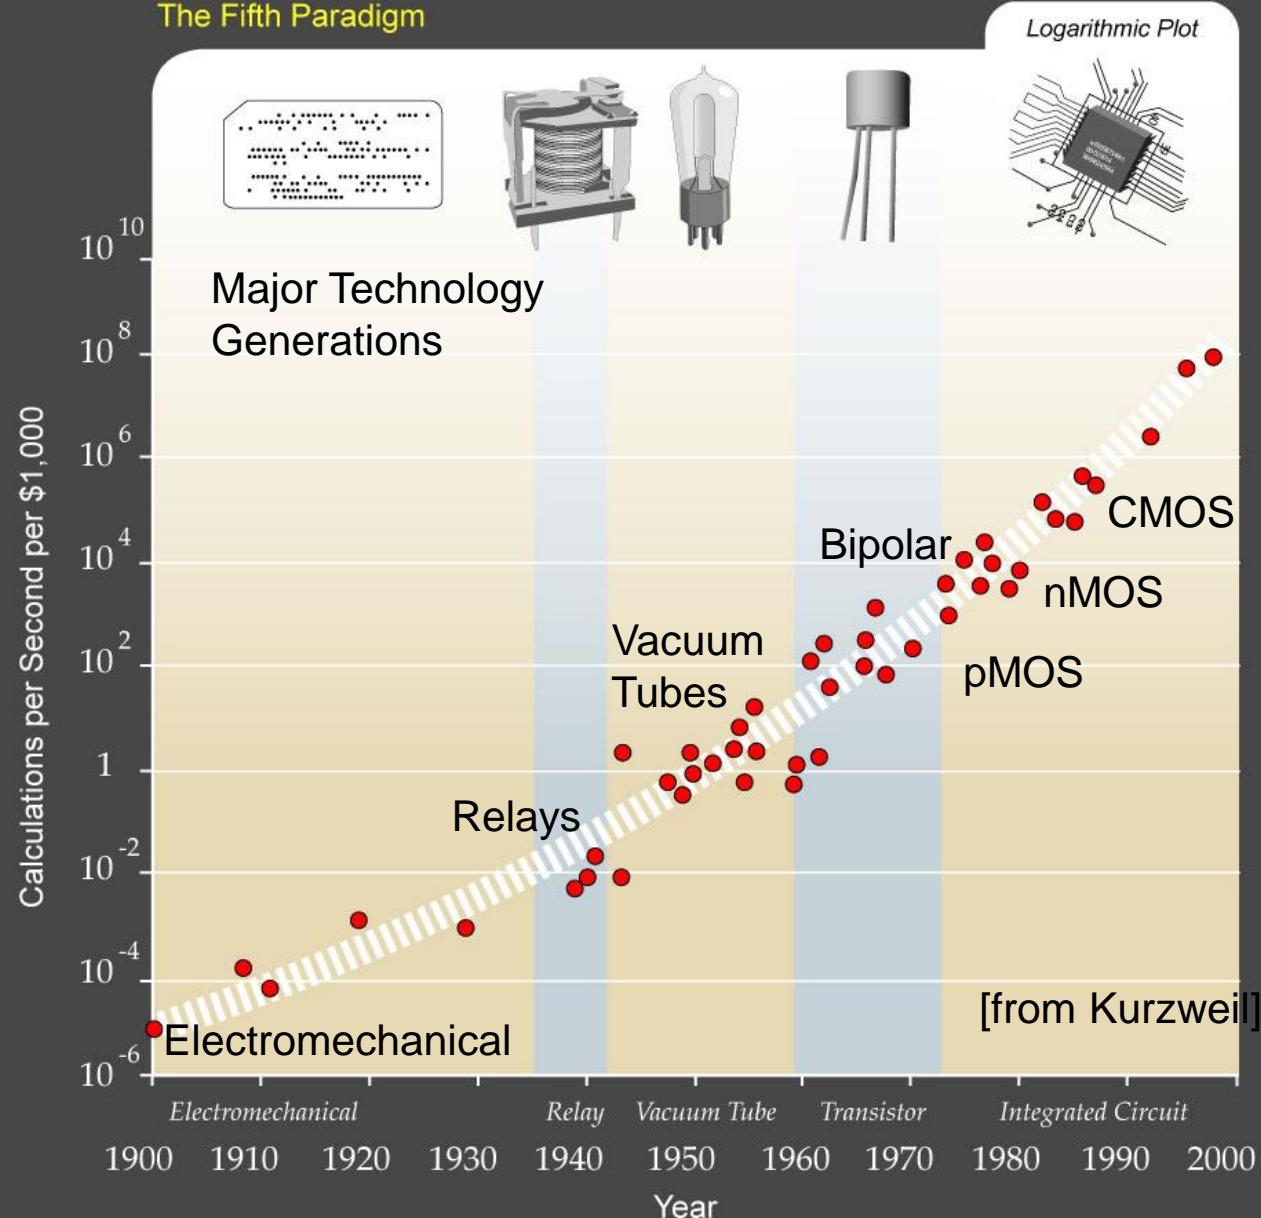
$$\text{Yield} = \frac{1}{(1 + (\text{Defects per area} \times \text{Die area}/2))^2}$$

- Nonlinear relation to area and defect rate
 - Wafer cost and area are fixed
 - Defect rate determined by manufacturing process
 - Die area determined by architecture and circuit design

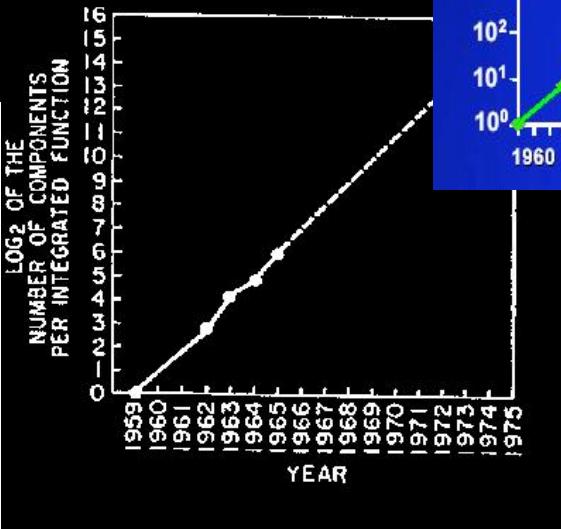
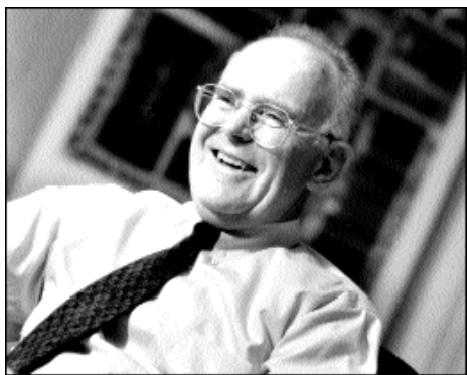
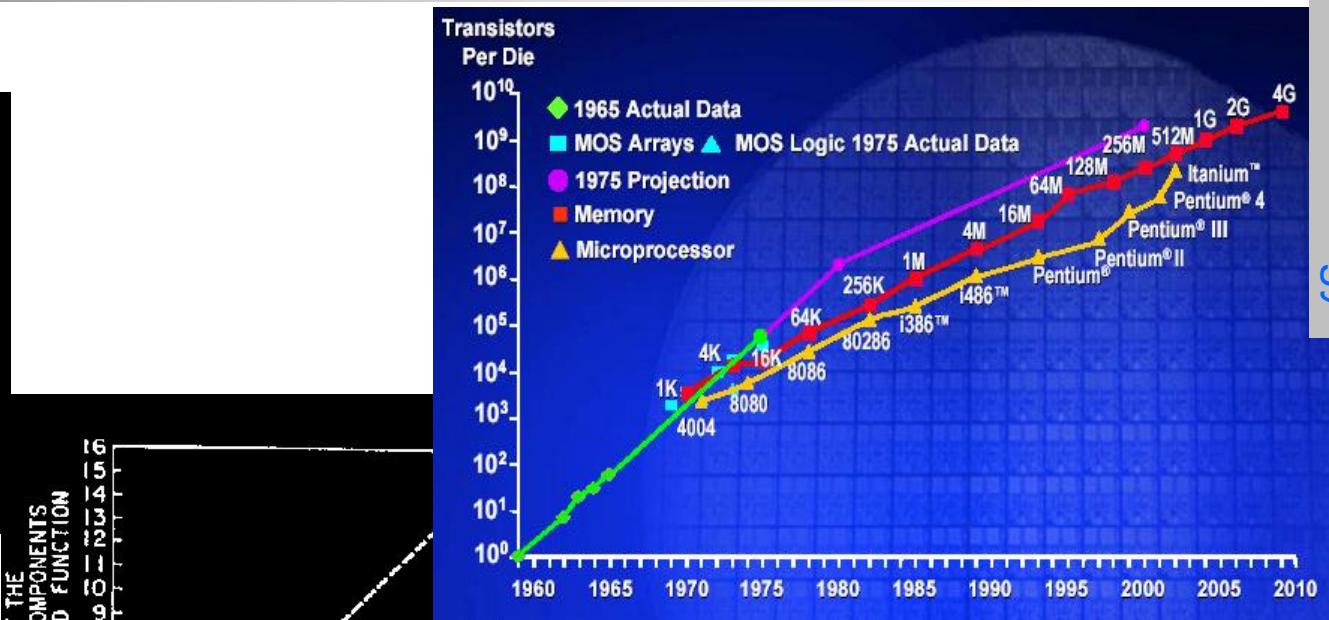
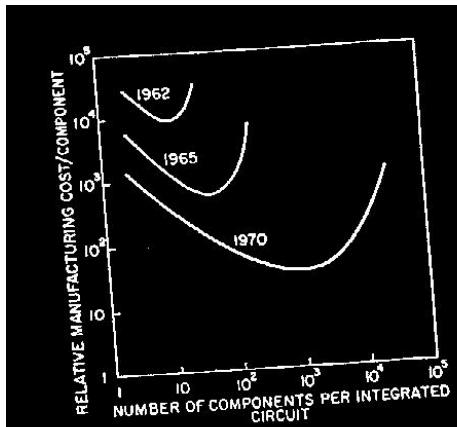
Moore's Law

The Fifth Paradigm

?



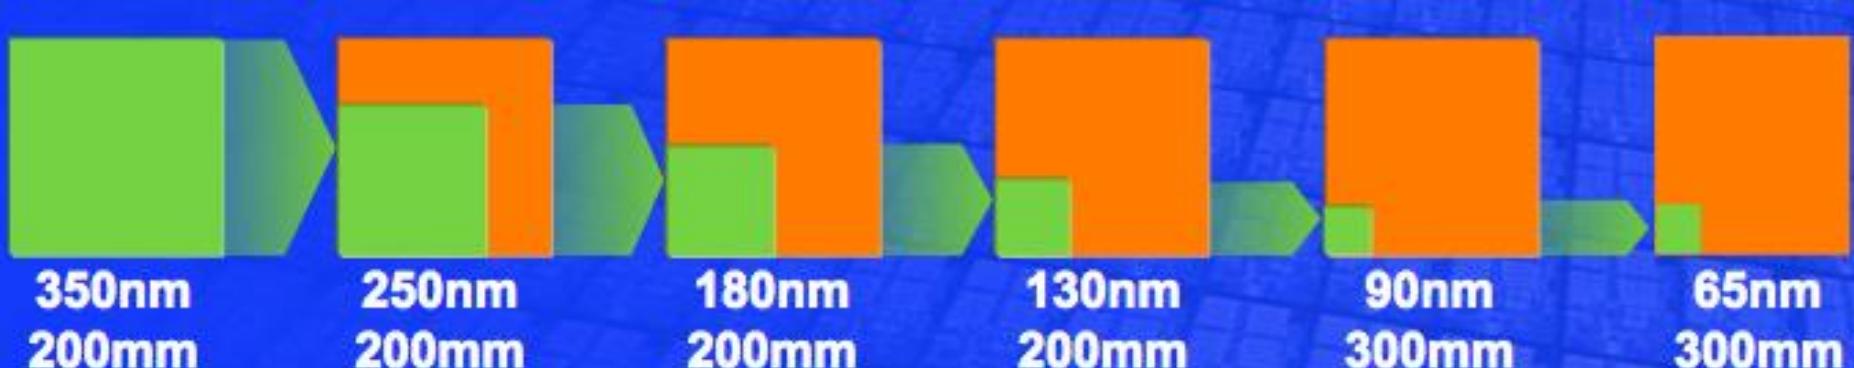
Moore's Law: 2X transistors / “year”



- “Cramming More Components onto Integrated Circuits”
 - Gordon Moore, Electronics, 1965
- # on transistors / cost-effective integrated circuit double every N months ($12 \leq N \leq 24$)



Main driver: device scaling ...



Twice the
circuitry in the
same space
(architectural
innovation)

OR

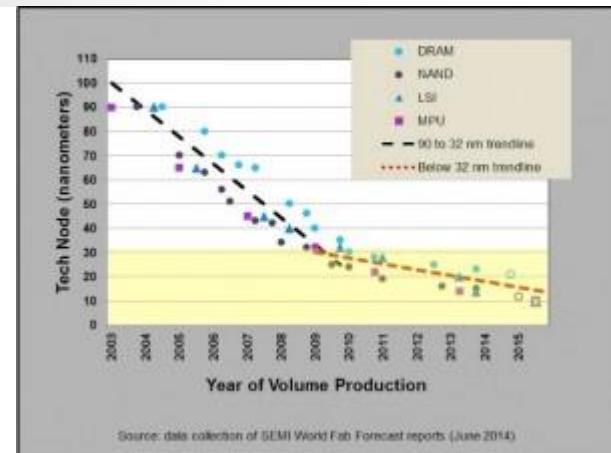
The same
circuitry in half
the space
(cost reduction)

Half the die size
for the same
capability than
in the prior
process

From: "Facing the Hot Chips Challenge Again", Bill Holt, Intel, presented at Hot Chips 17, 2005.

Trends in Technology

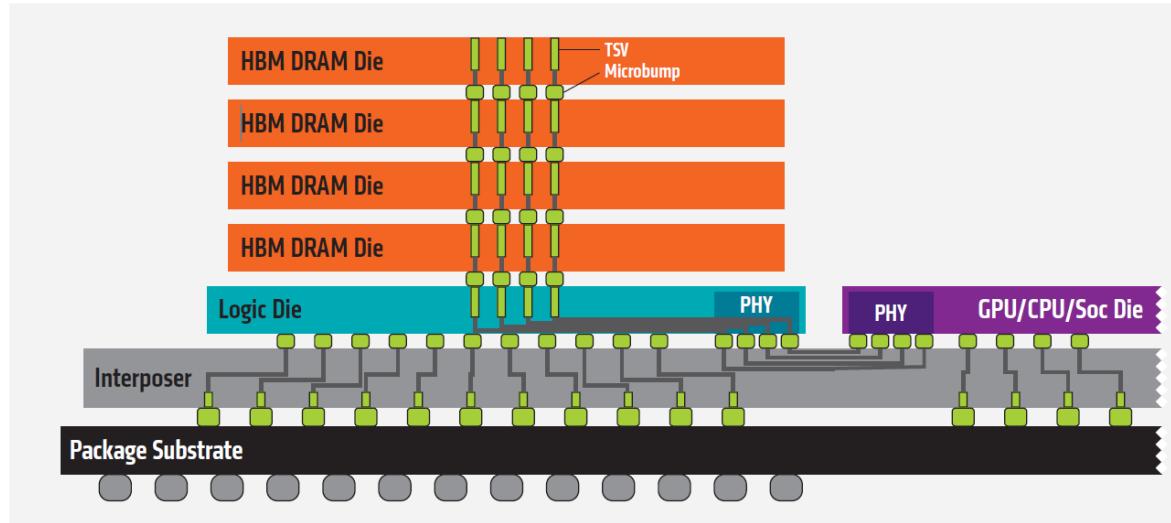
- Integrated circuit technology
 - Transistor density: 35%/year (slowing)
 - Die size: 10-20%/year
 - Integration overall: 40-55%/year (slowing)
 - Moore's Law is no more!
- DRAM capacity: was 25-40%/year (slowed dramatically now) (8Gbit in 2014, 16Gbit in 2019, 32Gbit ??)
- Flash capacity: 50-60%/year
 - 8-21X cheaper/bit than DRAM
- Magnetic disk technology: was 40%/year (~5% now)
 - 15-25X cheaper/bit then Flash
 - 300-500X cheaper/bit than DRAM



3D DRAM Stacking Technologies

Recent enabling technology: **3D stacking** of DRAM chips

- DRAMs connected via through-silicon-vias (TSVs) that run through the chips
 - TSVs provide highly parallel connection between logic layer and DRAMs
- Base layer of stack “**logic layer**” is memory controller, manages requests from processor
- Silicon “**interposer**” serves as high-BW interconnect between DRAM stack and processor



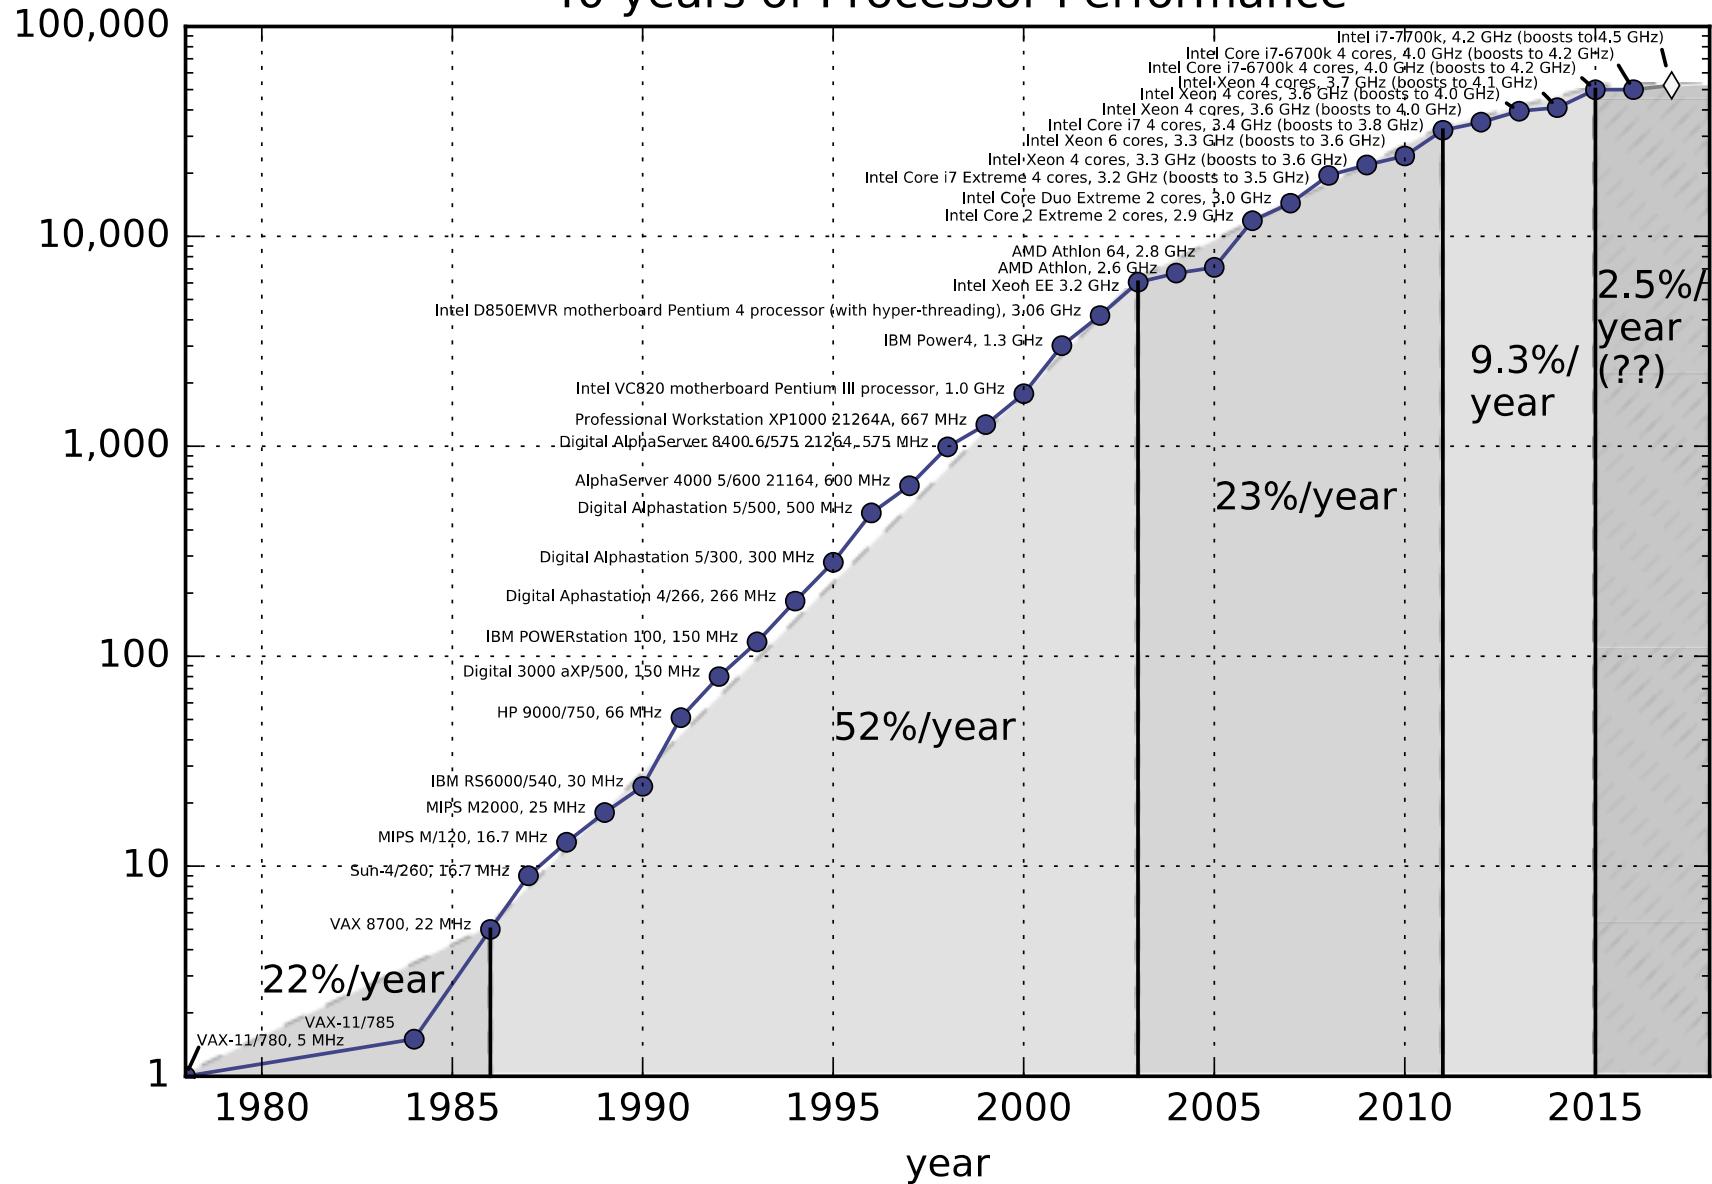
Technologies:

- Micron/Intel’s **Hybrid Memory Cube (HMC)**
- AMD’s **High-Bandwidth Memory (HBM)**: 1024 bit interface to stack

Single-Thread Processor Performance

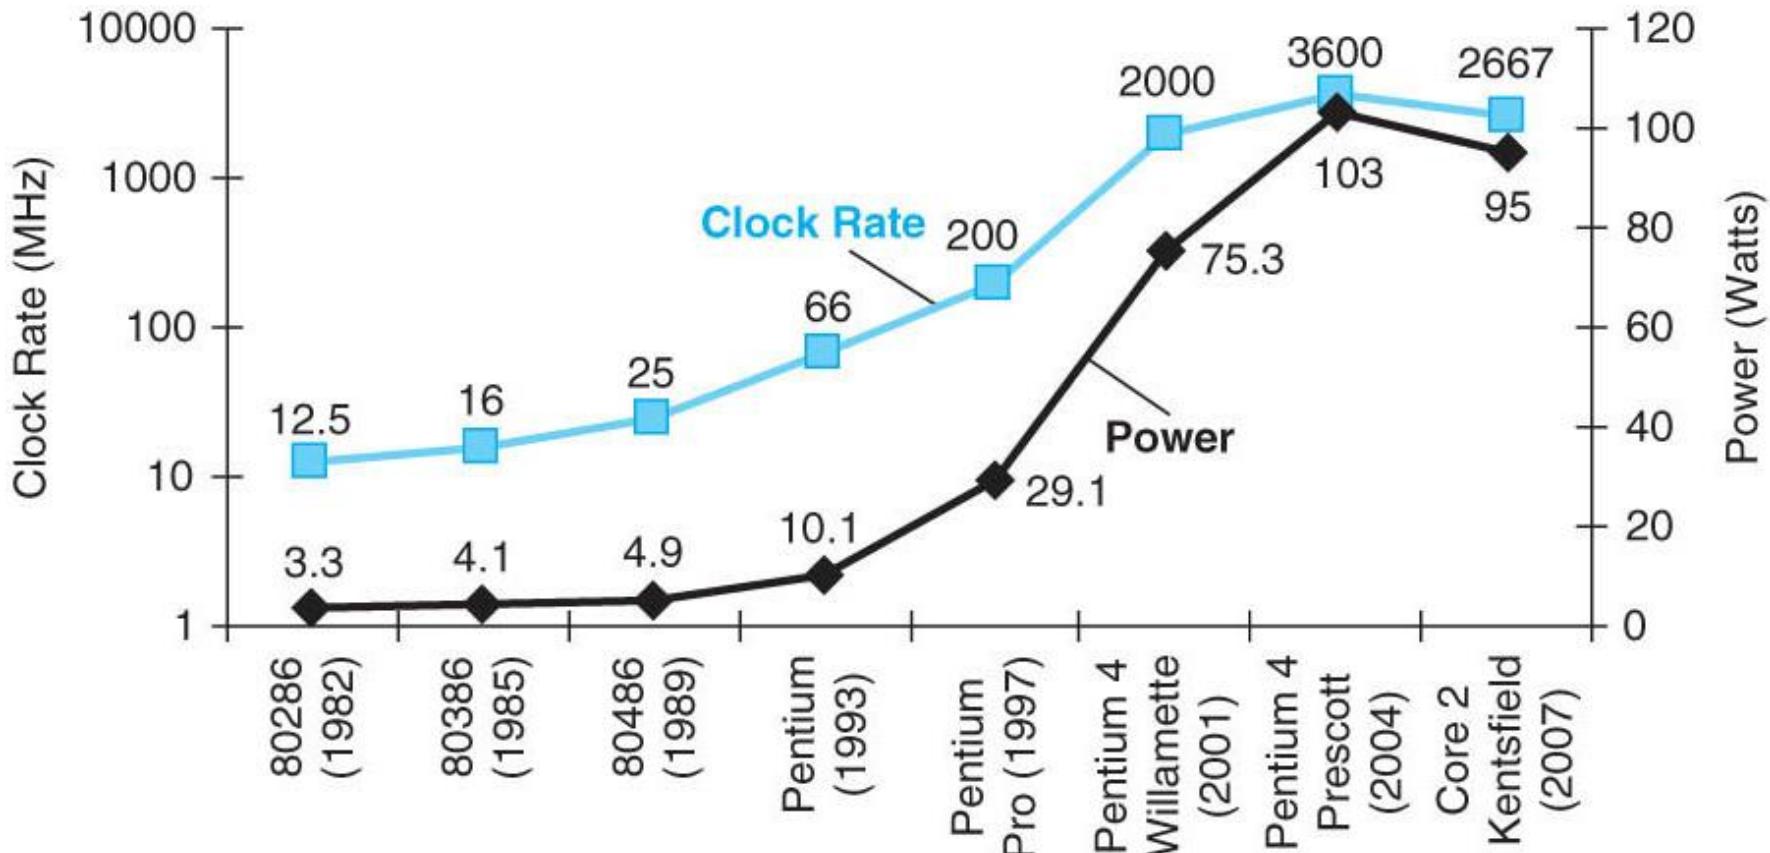
Performance vs. VAX-11/780

40 years of Processor Performance



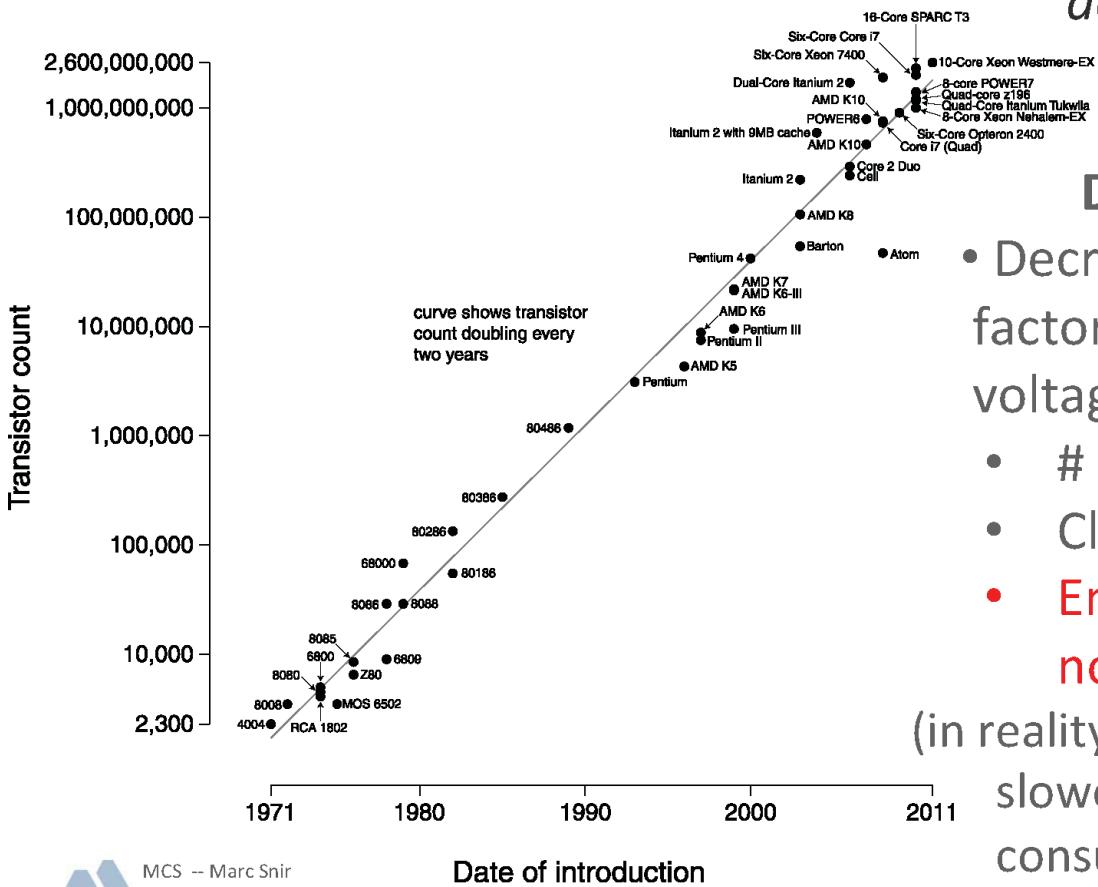
[Hennessy & Patterson, 2017]

Limits to Performance: Faster Means More Power



The CMOS Age: Moore's Law & Dennard Scaling

Microprocessor Transistor Counts 1971-2011 & Moore's Law



Moore's Law: The number of transistors per chip doubles every 2+ years

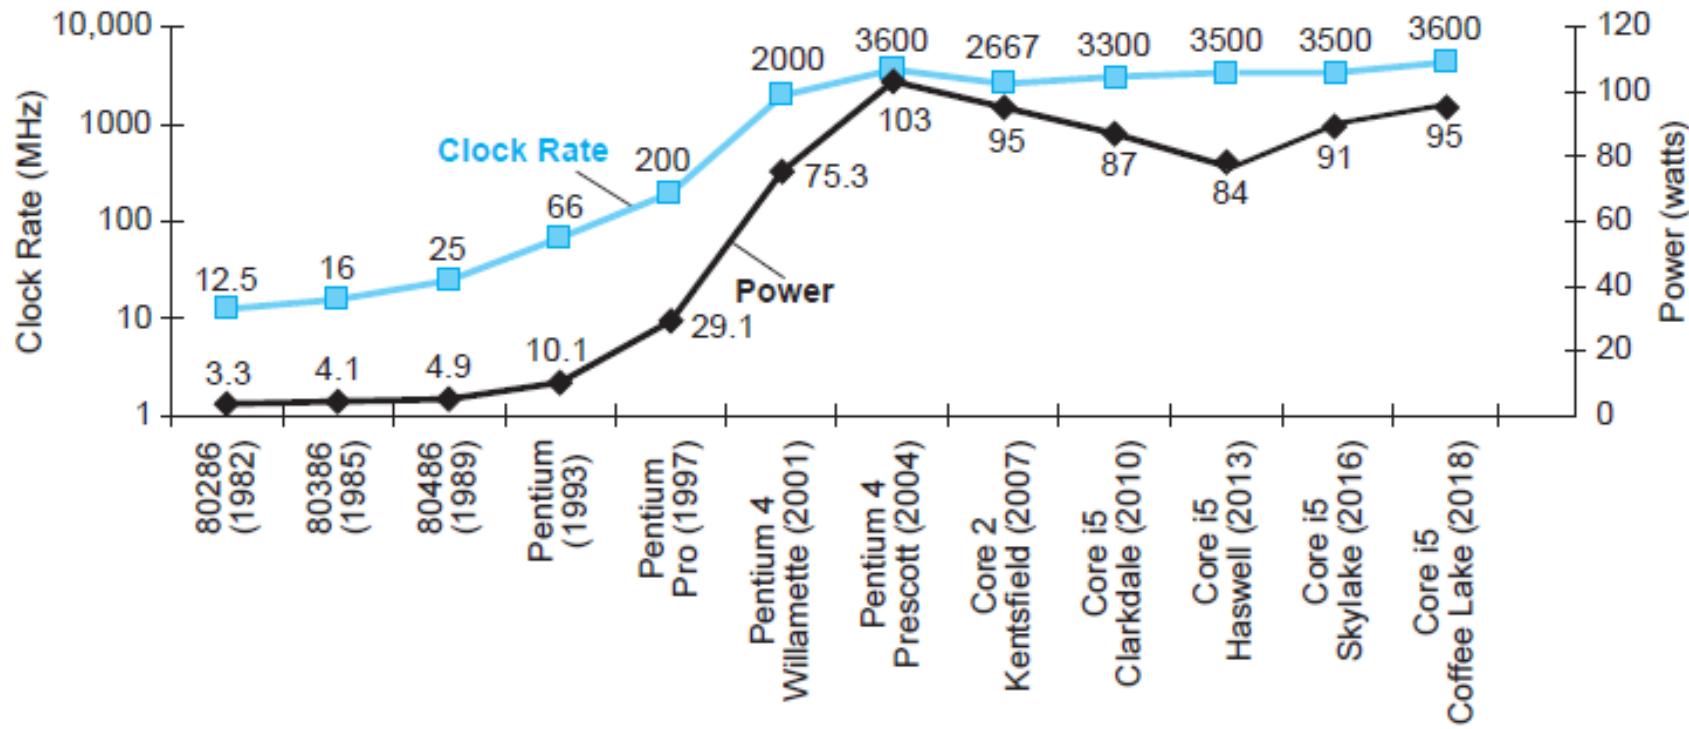
Dennard Scaling:

- Decrease feature size by a factor of λ and decrease voltage by a factor of λ ; then
 - # transistors increase by λ^2
 - Clock speed increases by λ
 - Energy consumption does not change

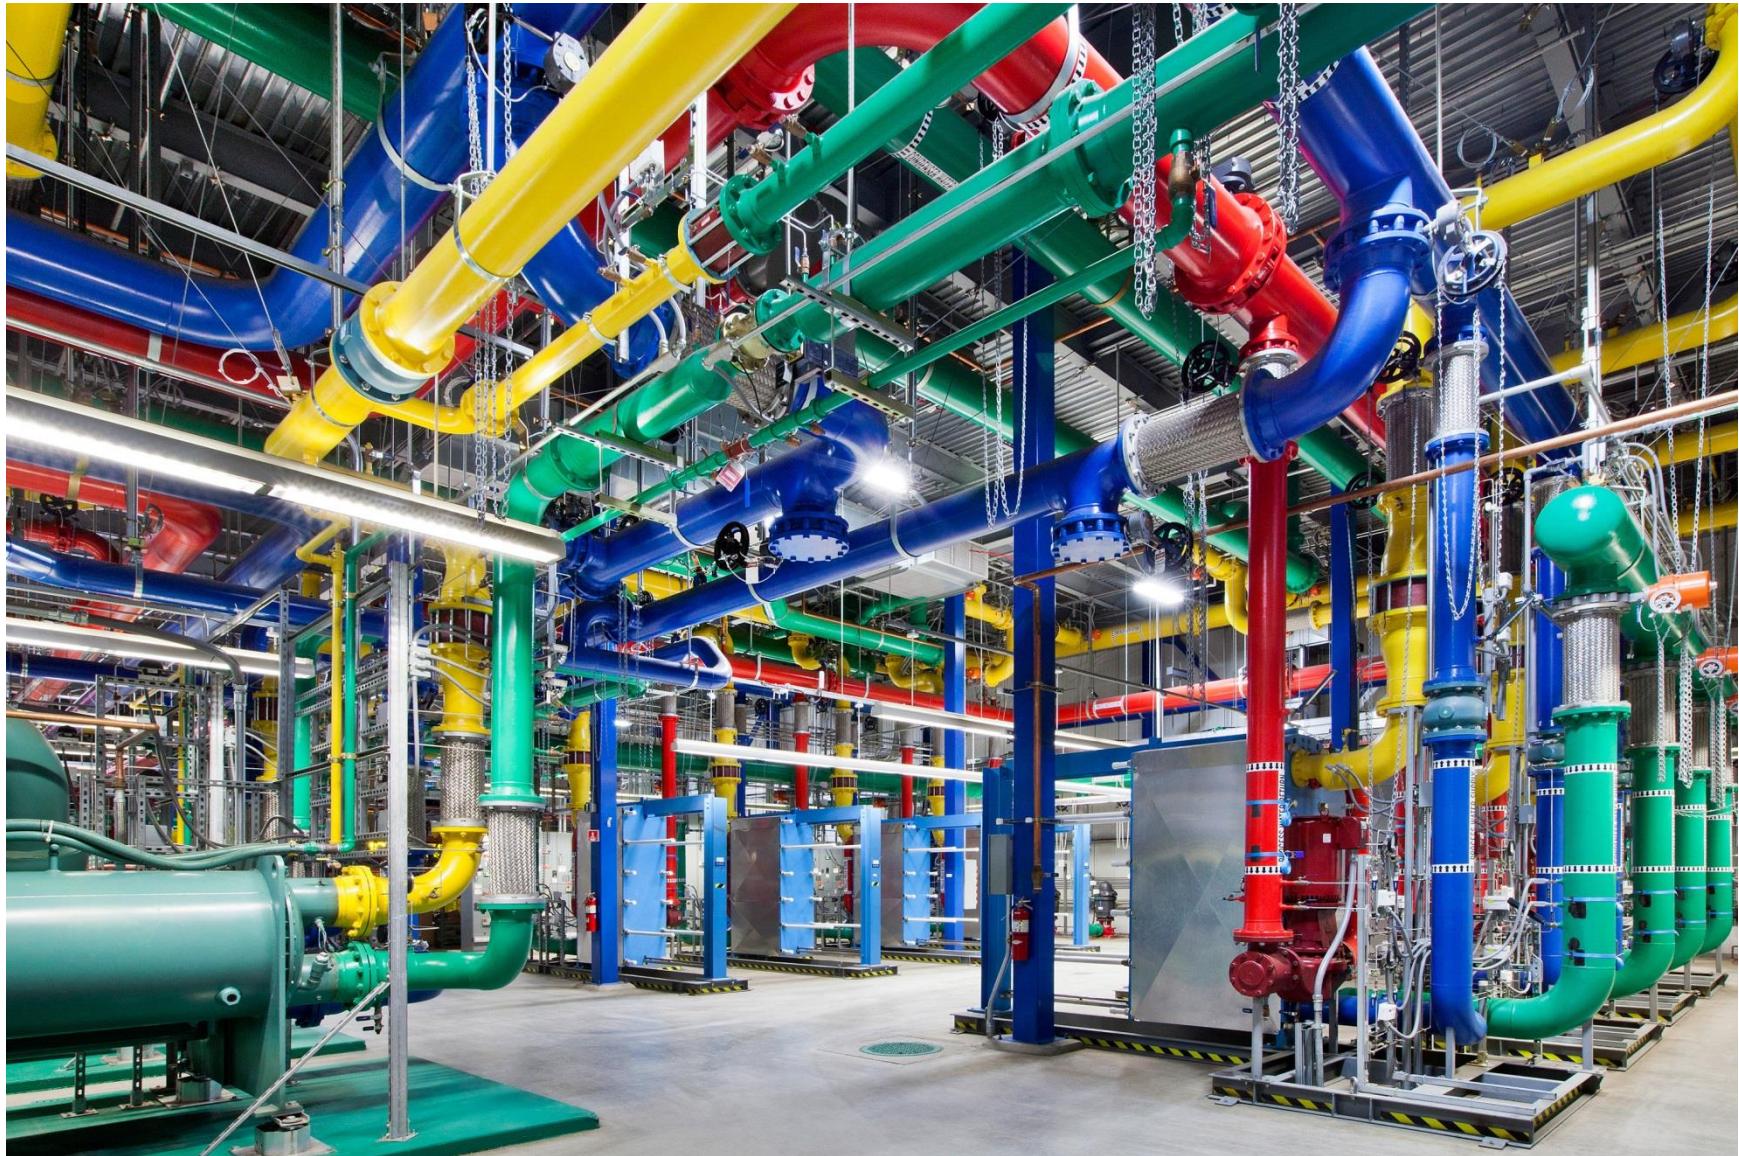
(in reality, voltage decrease was
slower; clock speed and energy
consumption increased faster)

The Power Wall

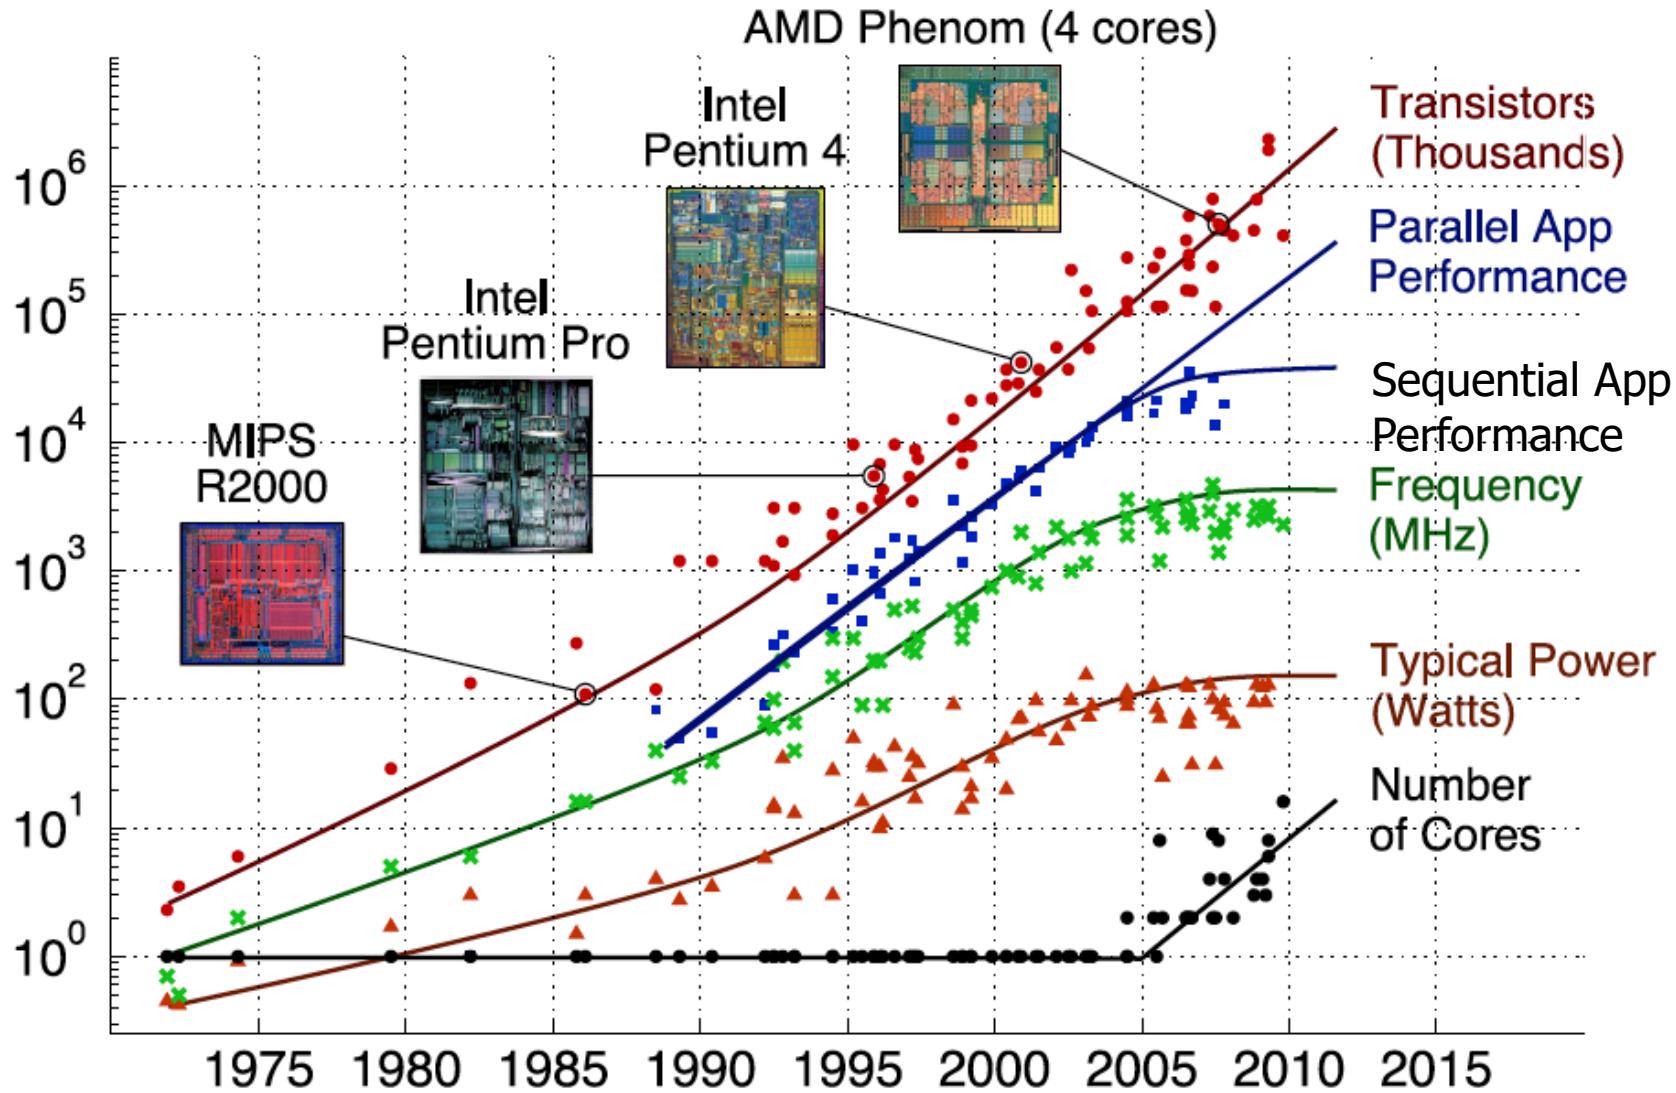
Why? We can no longer fully scale voltage Vdd because MOS transistor leakage current is no longer a negligible part of the power budget.



Water Cooling in a Google Data Center



Transition to Multicore



Data partially collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond

Multicore Machines Today

Examples from Apple's product line:



Mac Pro

28 Intel Xeon W cores



iMac Pro

18 Intel Xeon W cores



MacBook Pro Retina 15"

8 Intel Core i9 cores



iPad Pro

8 A12X cores

(4 fast +
4 low-power)



iPhone XS

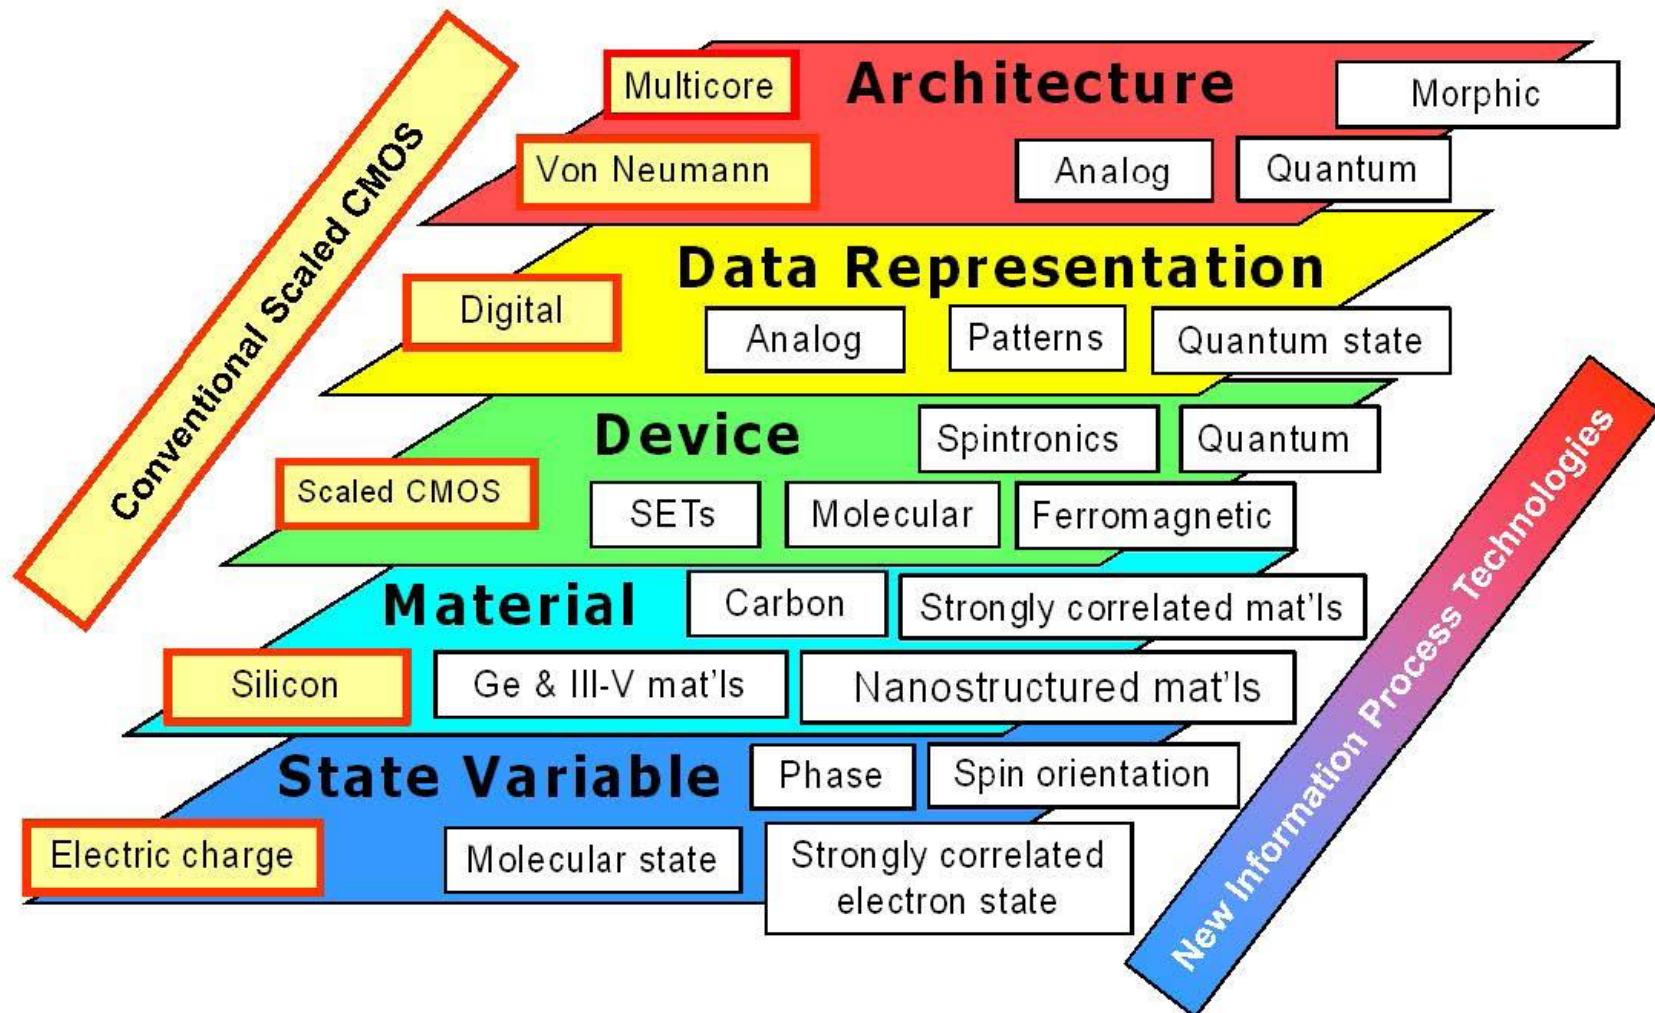
6 A12 cores

(2 fast +
4 low-power)

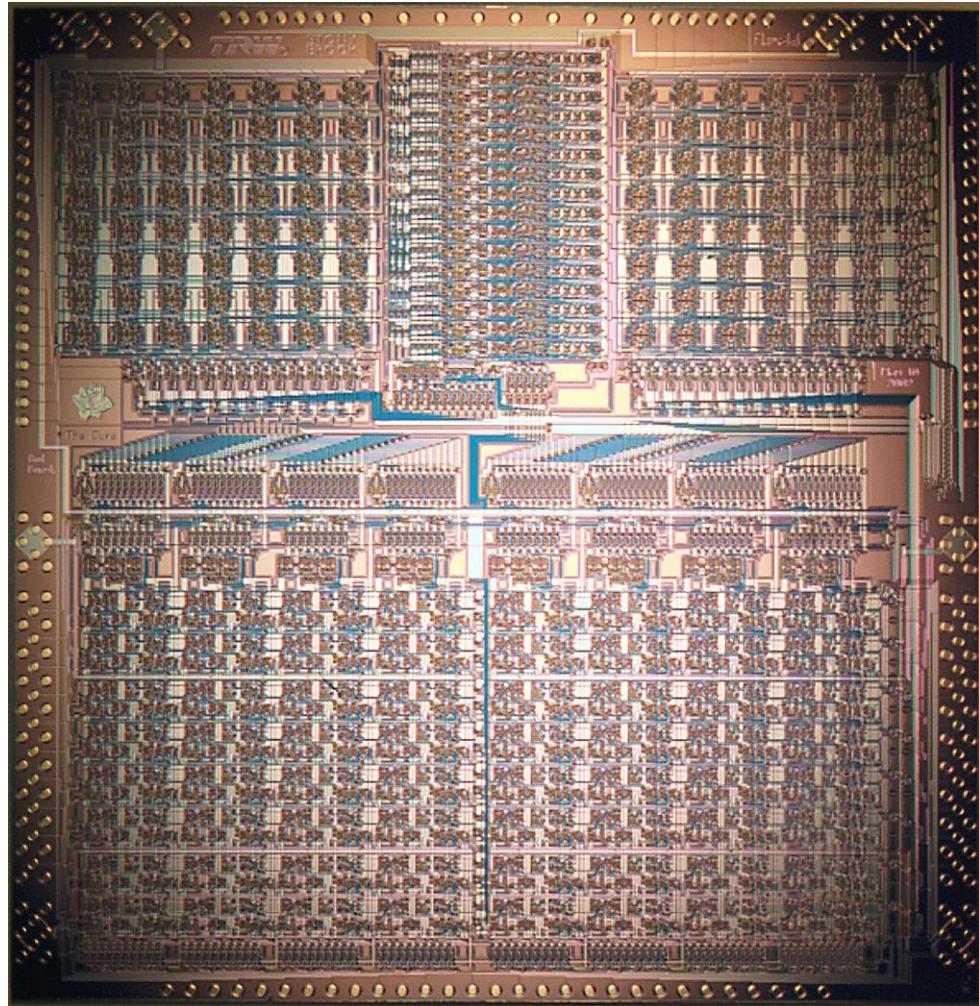


What's Next After CMOS?

A Taxonomy for Nano Information Processing Technologies



Superconductor Flux Quantum Computing: 20 GHz 8-bit RSFQ Microprocessor (2002)



- **Designed at Stony Brook University (ECE and Physics dept. teams) in collaboration with TRW (CA)**
- **Technology:** 1.75- μ m min. JJ size (TRW)
- **Circuit family:** RSFQ
- **Target Clock:** 20 GHz
- **JJ count:** 63,107
- **Power:** ~9.5 mW at 4.2K
- **Chip size:** 10.35 mm x 10.65 mm
- More about flux quantum computing at www.quantarctic.com

What About Bio/Neuro-Inspired Computing?

Human Brain

Laid out flat – 0.25 m² (not 3D) 2mm thick with 6 layers

wire diameter 10 nm, neuron diameter 4 microns

100 billion neurons (1×10^{11}),

0.15 quadrillion connections (1.5×10^{14})

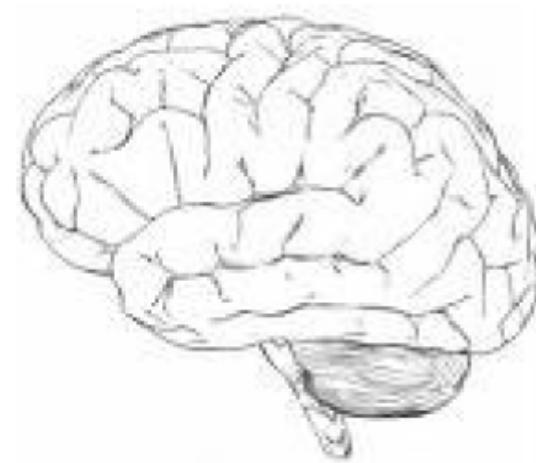
150,000 km of wiring

Wt 1.3 Kg

Frequency 30 hz

Power 25 W

System architects
eat your heart out



Continuous failure – 85,000 die/day, millions misfire/sec

Errors in calculations – make mistakes (examples visual illusions)

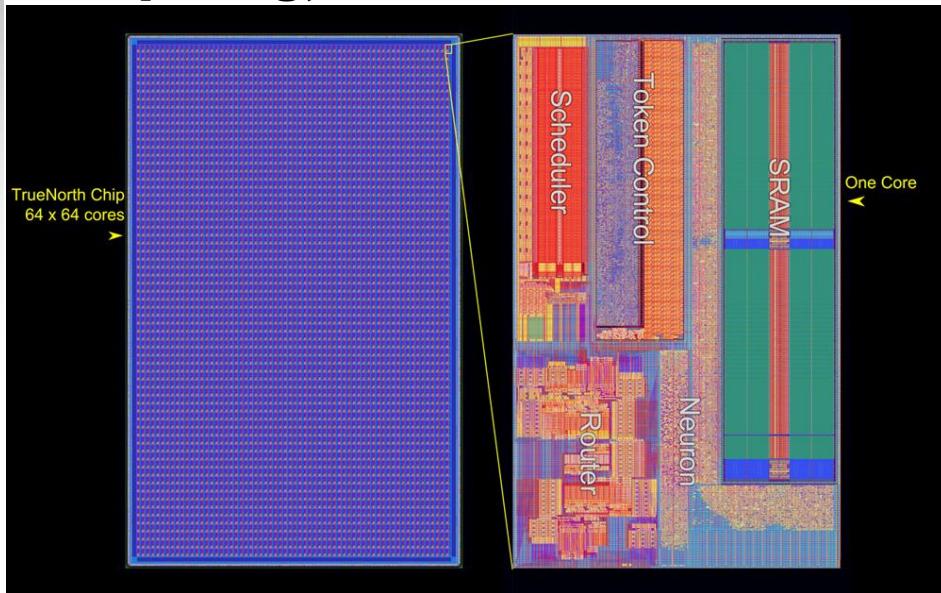
Self correction – sensor fusion, redundant calculation

Self healing – reprogram around physical damage



TrueNorth Neuromorphic Digital Chip (2014)

Low-power neuromorphic chip designed for applications in mobile sensors, cloud computing, and so on.

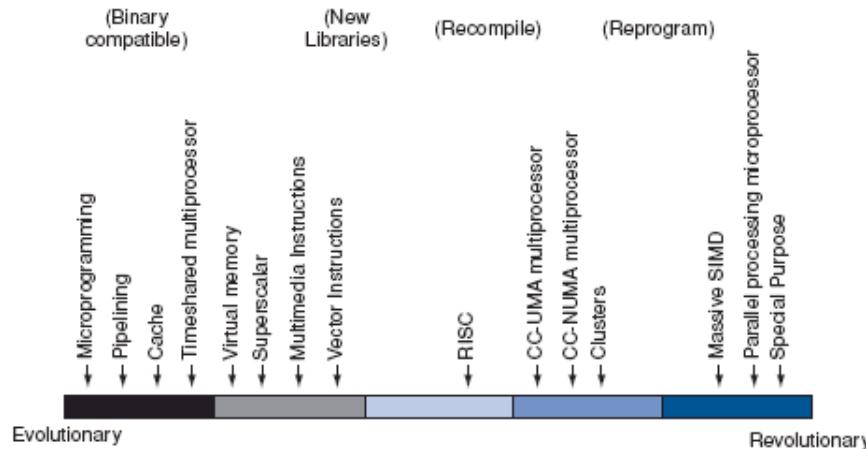


Each core has 256 inputs (axons), 256 outputs (neurons), a big bank of SRAM (which stores the data for each neuron), and a router that allows for any neuron to transmit to any axon up to 255 cores away.

- IBM, DARPA SyNAPSE
- 1 million neurons
- 256 million synapses
- 4096 neurosynaptic cores
- 5.4 billion transistors
- Asynchronous circuit design
- 28 nm Samsung Technology:
- 72 mW/op
- Largest current configuration:
 - 16 chips; 16 million neurons; 4 billion synapses
- Next configuration: 4,096 chips; 4 billion neurons; 1 trillion synapses
- Final configuration: 10 billion neurons; 100 trillion synapses

And in Conclusion ...

- Computer Architecture >> ISAs
- Computer architecture is shaped by technology and applications
 - History provides lessons for the future
- Computer Science at the crossroads from sequential to parallel computing
 - Salvation requires innovation in many fields, including computer architecture
- Evolution vs. Revolution
 - “More often the expense of innovation comes from being too disruptive to computer users”



- “Acceptance of hardware ideas requires acceptance by software people; therefore hardware people should learn about software. And if software people want good machines, they must learn more about hardware to be able to communicate with and thereby influence hardware engineers.”

Acknowledgements

- These slides contain material developed and copyright by:
 - Morgan Kauffmann (Elsevier, Inc.)
 - Arvind (MIT)
 - Krste Asanovic (MIT/UCB)
 - Joel Emer (Intel/MIT)
 - James Hoe (CMU)
 - John Kubiatowicz (UCB)
 - David Patterson (UCB)
 - Justin Hsia (UCB)
 - Mikhail Dorojevets (SBU)